

Joins and logic

Week 4

AEM 2850 / 5850 : R for Business Analytics
Cornell Dyson
Fall 2025

Acknowledgements: Grant McDermott, Jenny Bryan, R4DS (2e), Garrick Aden-Buie

Announcements

Reminders:

- Submit assignments via canvas / gradescope
 - Homework - Week 3 was due yesterday (Monday) at 11:59pm

Questions before we get started?

Plan for this week

Tuesday

Prologue

Joins

example-04-1

Thursday

Logic

- Boolean algebra
- Conditional transformations

example-04-2

Prologue

What sports do we watch?

Take a guess: what's the most popular spectator sport among classmates?

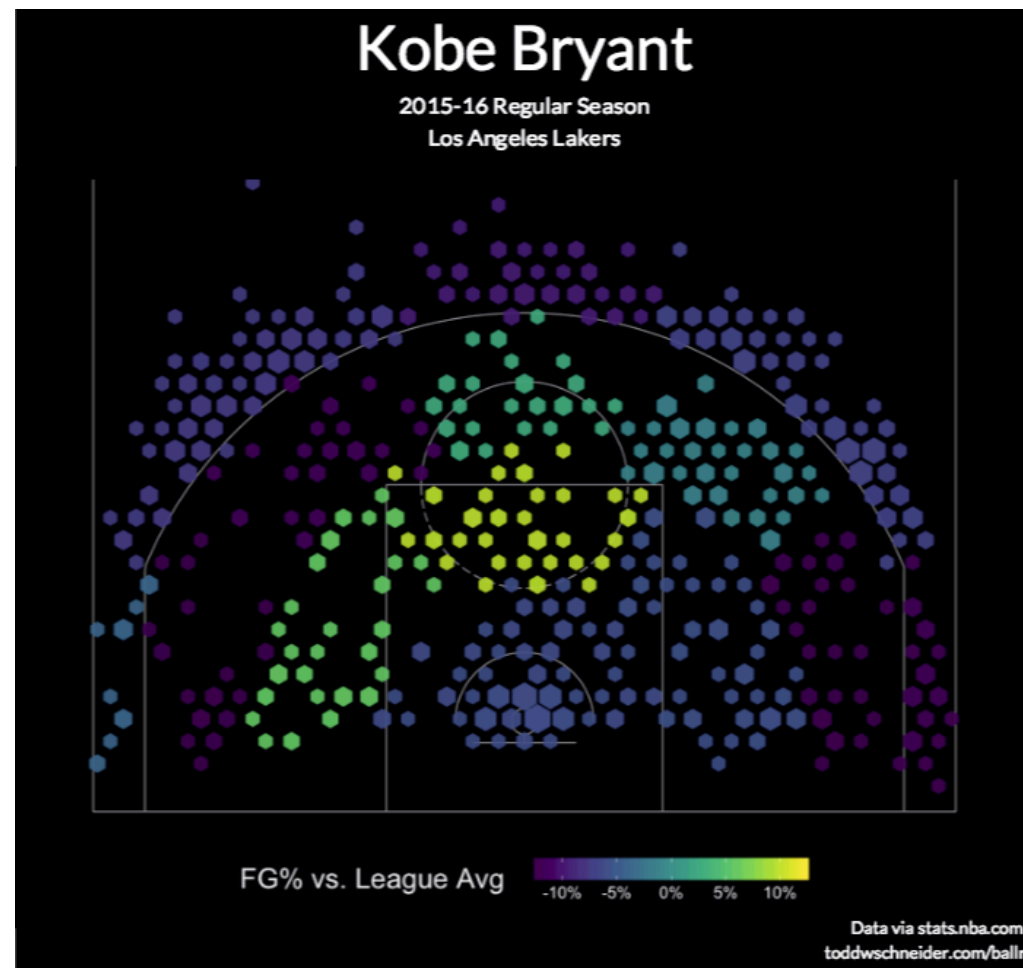
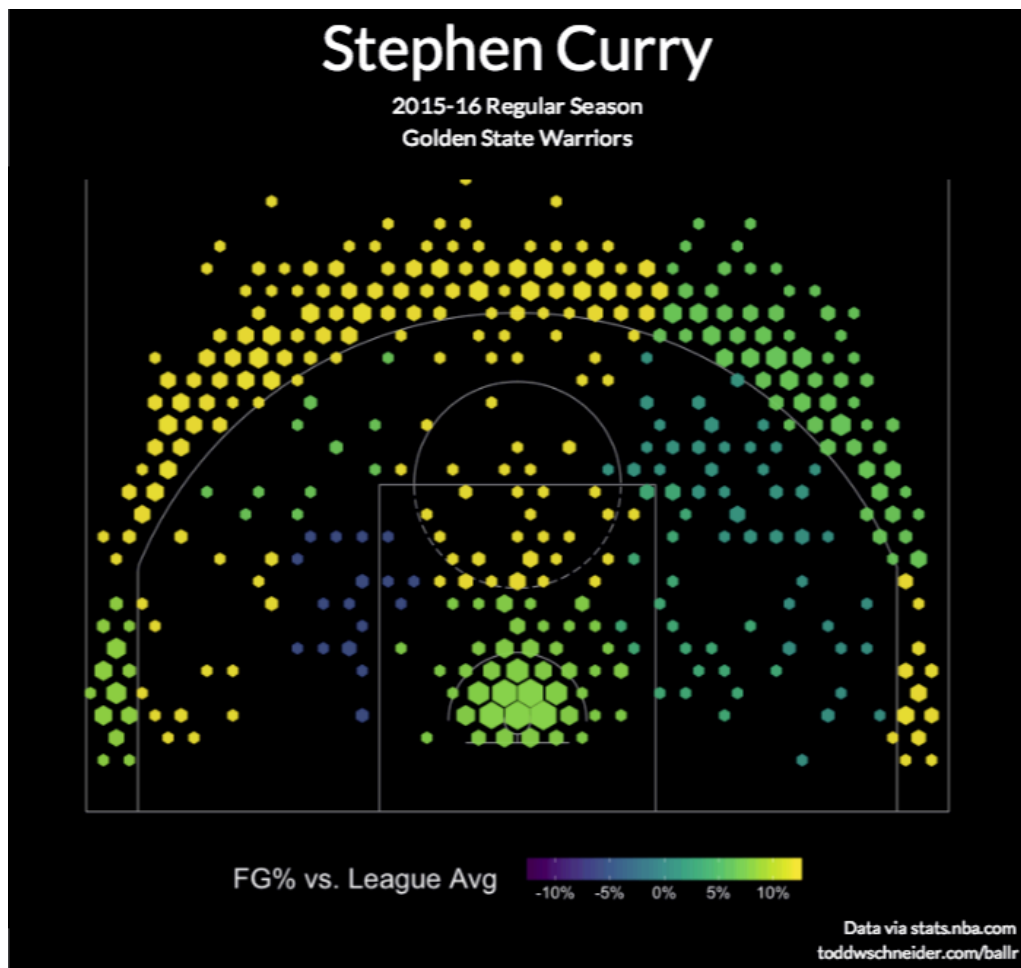
Here are the first 20 responses:

```
## [1] "soccer"      "baseball"    "baseball"    "badminton"   "football"
## [6] "baseball"    "volleyball"  "swimming"    "football"    "tennis"
## [11] "soccer"      "football"    "soccer"      "volleyball"  "basketball"
## [16] "hockey"      "tennis"      "volleyball"  "baseball"    "soccer"
```

Let's **count** and **arrange** to get the top 3:

```
## # A tibble: 3 × 2
##   sport      n
##   <chr>    <int>
## 1 basketball 14
## 2 baseball   12
## 3 soccer     12
```

R can be used for sports analytics, too!



Joins

Joins

Most data analyses require information contained in multiple data frames

We **join** them together to answer questions

Keys are the variables that connect a pair of data frames in a join

Join verbs from dplyr

1. **Mutating joins**: add new variables

- `left_join()`
- `right_join()`
- `inner_join()`
- `full_join()`

2. **Filtering joins**: filter observations

- `semi_join()`
- `anti_join()`

Join animations

Let's start by visualizing joins

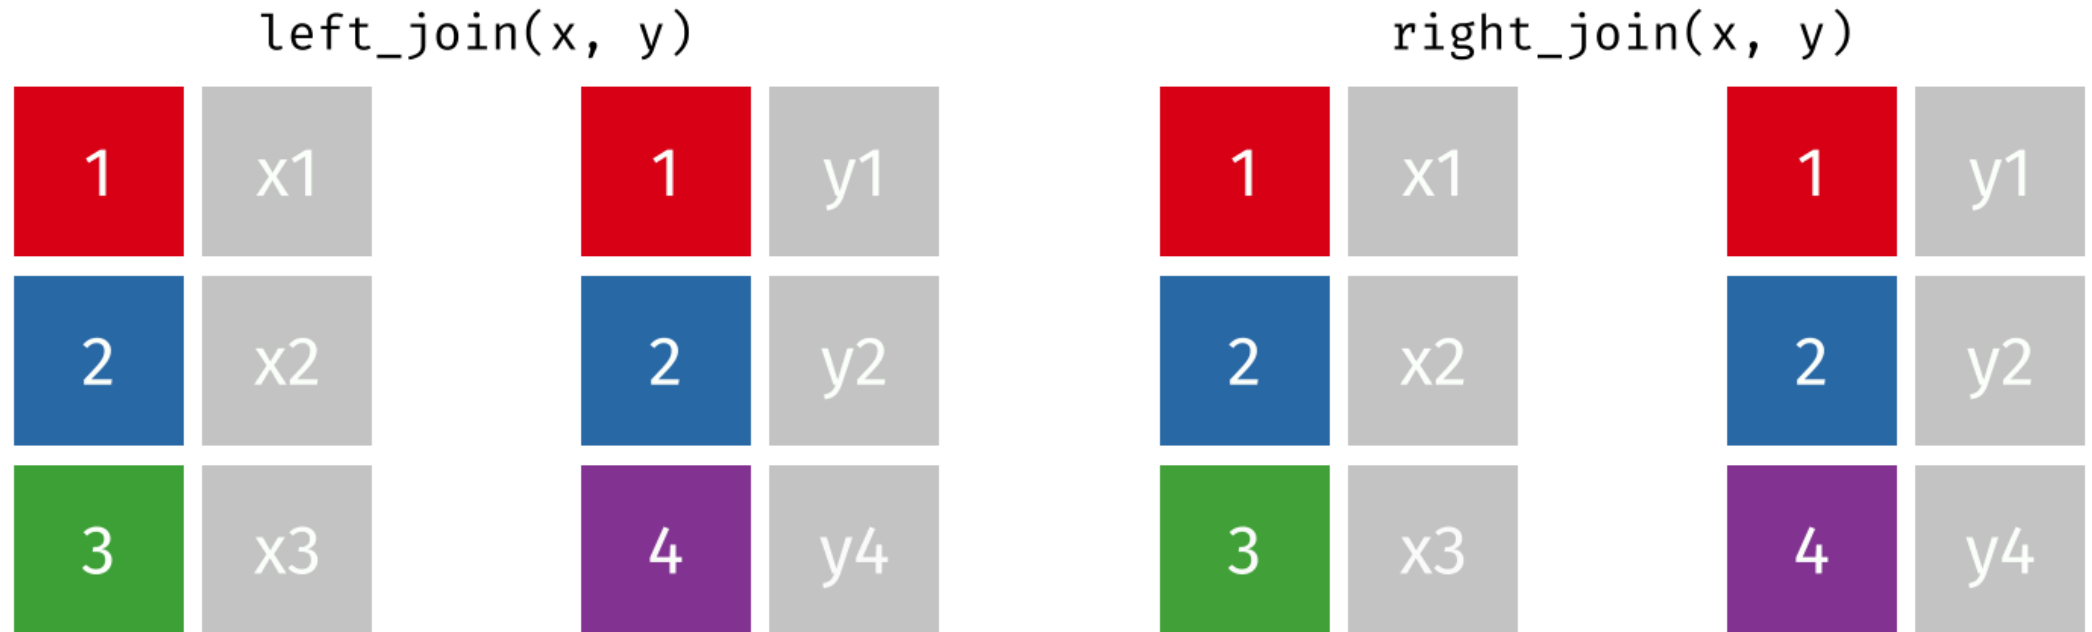
Here are two data frames we want to **join**

X		y	
1	x1	1	y1
2	x2	2	y2
3	x3	4	y4

Their **keys** are in color in the first column, and other data are in grey

Left join and right join

Left or right joins add variables to the left or right data frames



Multiple matches

With multiple matches between **x** and **y**, all combinations of matches are returned

left_join(x, y)

1	x1	1	y1
2	x2	2	y2
3	x3	4	y4
		2	y5

In this example, **x2** is duplicated to join one row in **x** to multiple rows in **y**

Inner join and full join

Inner joins return all rows in **x AND y**

`inner_join(x, y)`

1	x1	1	y1
2	x2	2	y2
3	x3	4	y4

Full joins return all rows in **x OR y**

`full_join(x, y)`

1	x1	1	y1
2	x2	2	y2
3	x3	4	y4

Semi join and anti join

Semi joins filter rows in **x** that match **y**

semi_join(x, y)			
1	x1	1	y1
2	x2	2	y2
3	x3	4	y4

Anti joins filter rows in **x** **not** in **y**

anti_join(x, y)			
1	x1	1	y1
2	x2	2	y2
3	x3	4	y4

example-04-1

Additional slides on joins for your reference

Joins

Let's learn these join commands using two small data frames

superheroes

```
## # A tibble: 7 × 3
##   name      alignment publisher
##   <chr>    <chr>      <chr>
## 1 Magneto  bad         Marvel
## 2 Storm    good        Marvel
## 3 Mystique bad         Marvel
## 4 Batman   good        DC
## 5 Joker    bad         DC
## 6 Catwoman bad         DC
## 7 Hellboy  good        Dark Horse Comics
```

publishers

```
## # A tibble: 3 × 2
##   publisher year_founded
##   <chr>      <int>
## 1 DC         1934
## 2 Marvel     1939
## 3 Image      1992
```

1) dplyr::left_join(x, y)

```
left_join(superheroes, publishers)
```

```
## Joining with `by = join_by(publisher)`
```

```
## # A tibble: 7 × 4
```

	name	alignment	publisher	year_founded
	<chr>	<chr>	<chr>	<int>
## 1	Magneto	bad	Marvel	1939
## 2	Storm	good	Marvel	1939
## 3	Mystique	bad	Marvel	1939
## 4	Batman	good	DC	1934
## 5	Joker	bad	DC	1934
## 6	Catwoman	bad	DC	1934
## 7	Hellboy	good	Dark Horse Comics	NA

`left_join` is a **mutating join**: it adds variables to `x`

`left_join` returns all rows from `x`

2) dplyr::right_join(x, y)

```
right_join(superheroes, publishers)
```

```
## Joining with `by = join_by(publisher)`
```

```
## # A tibble: 7 × 4
```

```
##   name      alignment publisher year_founded
##   <chr>      <chr>      <chr>      <int>
## 1 Magneto   bad        Marvel      1939
## 2 Storm     good        Marvel      1939
## 3 Mystique  bad        Marvel      1939
## 4 Batman    good        DC           1934
## 5 Joker     bad        DC           1934
## 6 Catwoman  bad        DC           1934
## 7 <NA>      <NA>      Image       1992
```

`right_join` is a **mutating join**: it adds variables to `y`

`right_join` returns all rows from `y`

3) dplyr::inner_join(x, y)

```
inner_join(superheroes, publishers)
```

```
## Joining with `by = join_by(publisher)`
```

```
## # A tibble: 6 × 4
```

```
##   name      alignment publisher year_founded
##   <chr>     <chr>      <chr>      <int>
## 1 Magneto   bad         Marvel      1939
## 2 Storm     good        Marvel      1939
## 3 Mystique  bad         Marvel      1939
## 4 Batman    good         DC           1934
## 5 Joker     bad          DC           1934
## 6 Catwoman  bad          DC           1934
```

How is `inner_join` different from `left_join` and `right_join`?

`inner_join` returns all rows in `x` **AND** `y`

4) dplyr::full_join(x, y)

```
full_join(superheroes, publishers) # how many rows do you think this will produce?
```

```
## Joining with `by = join_by(publisher)`
```

```
## # A tibble: 8 × 4
```

	name	alignment	publisher	year_founded
	<chr>	<chr>	<chr>	<int>
## 1	Magneto	bad	Marvel	1939
## 2	Storm	good	Marvel	1939
## 3	Mystique	bad	Marvel	1939
## 4	Batman	good	DC	1934
## 5	Joker	bad	DC	1934
## 6	Catwoman	bad	DC	1934
## 7	Hellboy	good	Dark Horse Comics	NA
## 8	<NA>	<NA>	Image	1992

full_join returns all rows in x **OR** y

5) dplyr::semi_join(x, y)

```
superheroes
```

```
## # A tibble: 7 × 3
##   name      alignment publisher
##   <chr>    <chr>      <chr>
## 1 Magneto  bad        Marvel
## 2 Storm    good        Marvel
## 3 Mystique bad        Marvel
## 4 Batman   good        DC
## 5 Joker    bad        DC
## 6 Catwoman bad        DC
## 7 Hellboy  good        Dark Horse Comics
```

```
semi_join(superheroes, publishers)
```

```
## Joining with `by = join_by(publisher)`

## # A tibble: 6 × 3
##   name      alignment publisher
##   <chr>    <chr>      <chr>
## 1 Magneto  bad        Marvel
## 2 Storm    good        Marvel
## 3 Mystique bad        Marvel
## 4 Batman   good        DC
## 5 Joker    bad        DC
## 6 Catwoman bad        DC
```

`semi_join` is a **filtering join**: it keeps observations in `x` that have a match in `y`

Note that the variables do not change

6) dplyr::anti_join(x, y)

```
superheroes
```

```
## # A tibble: 7 × 3
##   name      alignment publisher
##   <chr>    <chr>      <chr>
## 1 Magneto  bad         Marvel
## 2 Storm    good        Marvel
## 3 Mystique bad         Marvel
## 4 Batman   good        DC
## 5 Joker    bad         DC
## 6 Catwoman bad         DC
## 7 Hellboy  good        Dark Horse Comics
```

```
anti_join(superheroes, publishers)
```

```
## Joining with `by = join_by(publisher)`

## # A tibble: 1 × 3
##   name      alignment publisher
##   <chr>    <chr>      <chr>
## 1 Hellboy  good        Dark Horse Comics
```

`anti_join` is a **filtering join**: it keeps obs. in `x` that **DO NOT** have a match in `y`

Note that the variables do not change

Key variables

How do `dplyr` join commands know what variables to use as **keys**?

By default, `*_join()` uses all variables that are common across `x` and `y`

```
intersect(names(superheroes), names(publishers)) # variable used for matching before
```

```
## [1] "publisher"
```

Or, we can specify what to join by: `*_join(..., by = join_by(publisher))`

Note: before `dplyr` 1.1.0, the syntax was: `*_join(..., by = "publisher")`

Exploring keys

```
library(nycflights13) # let's explore keys using the nycflights13 data
flights |> print(n = 8) # print the first 8 rows of flights
```

```
## # A tibble: 336,776 × 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517           515           2     830           819
## 2  2013     1     1     533           529           4     850           830
## 3  2013     1     1     542           540           2     923           850
## 4  2013     1     1     544           545          -1    1004          1022
## 5  2013     1     1     554           600          -6     812           837
## 6  2013     1     1     554           558          -4     740           728
## 7  2013     1     1     555           600          -5     913           854
## 8  2013     1     1     557           600          -3     709           723
## # i 336,768 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Exploring keys

```
planes # print the first 10 rows of planes
```

```
## # A tibble: 3,322 × 9
##   tailnum year type      manufacturer model engines seats speed engine
##   <chr>   <int> <chr>      <chr>          <chr>   <int> <int> <int> <chr>
## 1 N10156  2004 Fixed wing multi... EMBRAER      EMB-...     2    55    NA Turbo...
## 2 N102UW  1998 Fixed wing multi... AIRBUS INDU... A320...     2   182    NA Turbo...
## 3 N103US  1999 Fixed wing multi... AIRBUS INDU... A320...     2   182    NA Turbo...
## 4 N104UW  1999 Fixed wing multi... AIRBUS INDU... A320...     2   182    NA Turbo...
## 5 N10575  2002 Fixed wing multi... EMBRAER      EMB-...     2    55    NA Turbo...
## 6 N105UW  1999 Fixed wing multi... AIRBUS INDU... A320...     2   182    NA Turbo...
## 7 N107US  1999 Fixed wing multi... AIRBUS INDU... A320...     2   182    NA Turbo...
## 8 N108UW  1999 Fixed wing multi... AIRBUS INDU... A320...     2   182    NA Turbo...
## 9 N109UW  1999 Fixed wing multi... AIRBUS INDU... A320...     2   182    NA Turbo...
## 10 N110UW  1999 Fixed wing multi... AIRBUS INDU... A320...     2   182    NA Turbo...
## # i 3,312 more rows
```

Let's perform a left join on flights and planes

```
left_join(flights, planes) |>  
  select(year:dep_time, arr_time, carrier:tailnum, type, model) |> # keep text to one slide  
  print(n = 5) # just to save vertical space on the slide
```

```
## Joining with `by = join_by(year, tailnum)`
```

```
## # A tibble: 336,776 × 10
```

```
##   year month   day dep_time arr_time carrier flight tailnum type  model  
##   <int> <int> <int>   <int>   <int> <chr>   <int> <chr>   <chr> <chr>  
## 1  2013     1     1     517     830 UA      1545 N14228 <NA> <NA>  
## 2  2013     1     1     533     850 UA      1714 N24211 <NA> <NA>  
## 3  2013     1     1     542     923 AA      1141 N619AA <NA> <NA>  
## 4  2013     1     1     544    1004 B6       725 N804JB <NA> <NA>  
## 5  2013     1     1     554     812 DL       461 N668DN <NA> <NA>
```

```
## # i 336,771 more rows
```

Uh-oh! What's up with **type** and **model**?

Uh-oh!

As before, `dplyr` guessed which columns to join on

It uses columns with the same name:

```
## Joining, by = c("year", "tailnum")
```

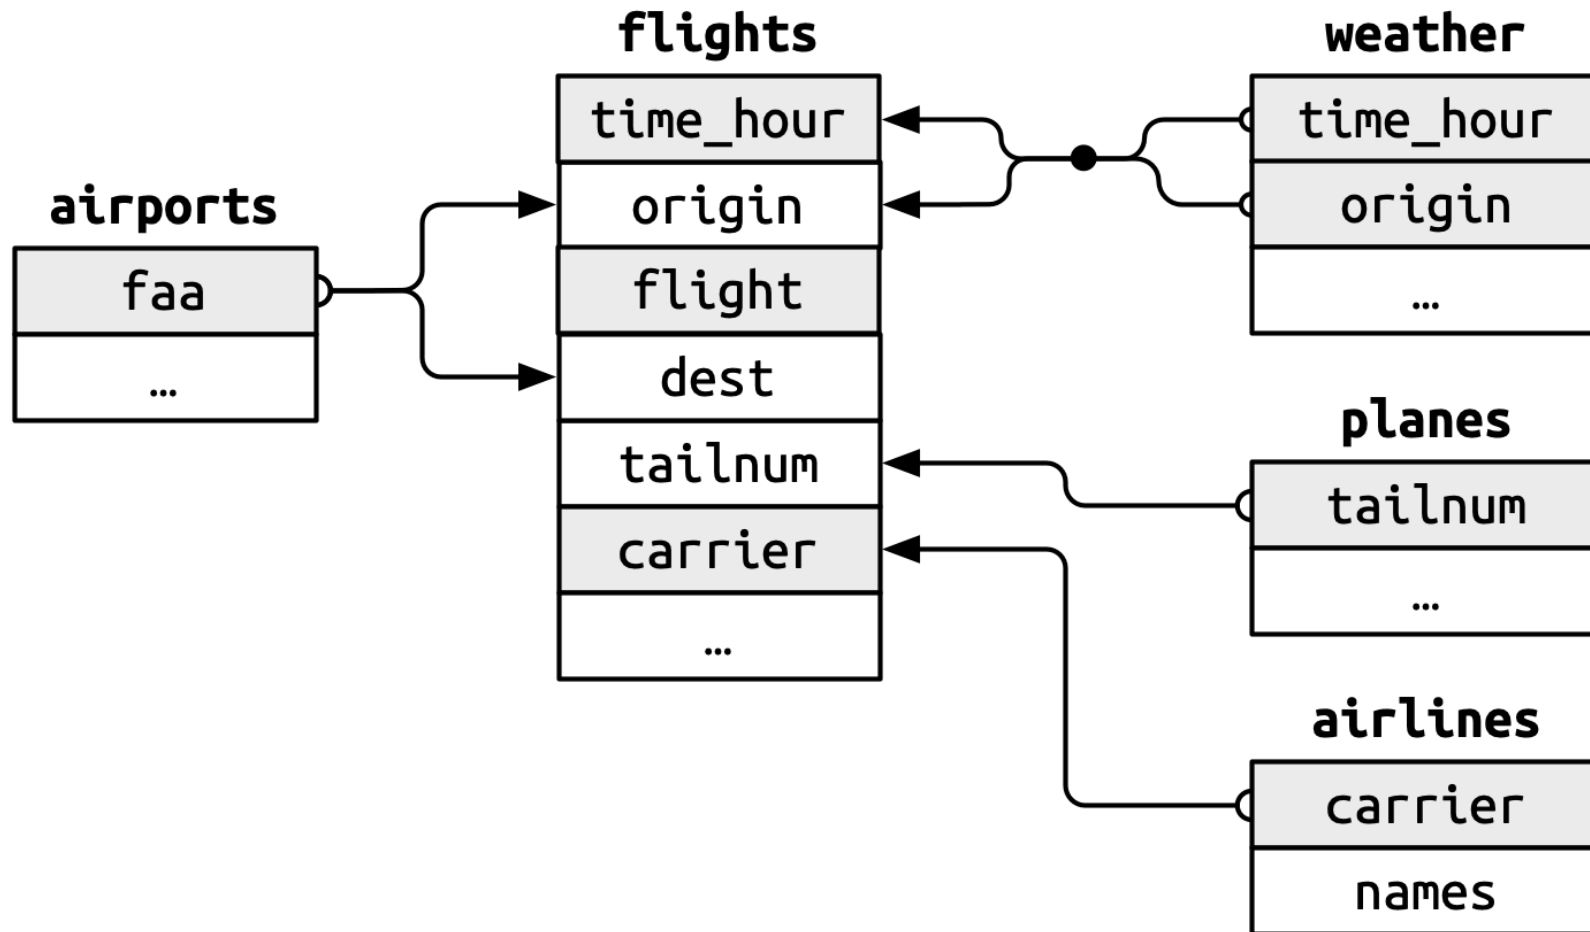
Does anyone see a potential problem here?

The variable `year` does not have a consistent meaning across the datasets

In `flights` it refers to the *year of flight*, in `planes` it refers to *year of construction*

Luckily we can avoid this by using the argument `by = join_by(...)`

What should we join flights and planes by?



Specifying join keys

We just need to be explicit in the join call by using the **by** argument

```
left_join(flights,
          planes |> rename(year_built = year), # not necessary w/ below line, but helpful
          by = join_by(tailnum) # be specific about the joining column
        ) |>
  select(year, month:dep_time, carrier, flight, tailnum, year_built, type, model) |>
  print(n = 5) # just to save vertical space on the slide
```

```
## # A tibble: 336,776 × 10
##   year month   day dep_time carrier flight tailnum year_built type      model
##   <int> <int> <int>   <int> <chr>   <int> <chr>      <int> <chr>    <chr>
## 1  2013     1     1     517 UA       1545 N14228     1999 Fixed wing... 737-...
## 2  2013     1     1     533 UA       1714 N24211     1998 Fixed wing... 737-...
## 3  2013     1     1     542 AA       1141 N619AA     1990 Fixed wing... 757-...
## 4  2013     1     1     544 B6        725 N804JB     2012 Fixed wing... A320...
## 5  2013     1     1     554 DL        461 N668DN     1991 Fixed wing... 757-...
## # i 336,771 more rows
```

Specifying join keys

What happens if we don't rename `year` before this join?

```
left_join(flights,  
          planes, # not renaming "year" to "year_built" this time  
          by = join_by(tailnum)  
        ) |>  
  select(contains("year"), month:dep_time, arr_time, carrier, flight, tailnum, type, model) |>  
  print(n = 4) # just to save vertical space on the slide
```

```
## # A tibble: 336,776 × 11  
##   year.x year.y month   day dep_time arr_time carrier flight tailnum type  model  
##   <int> <int> <int> <int>   <int>   <int> <chr>   <int> <chr>  <chr> <chr>  
## 1  2013   1999     1     1     517     830 UA      1545 N14228 Fixe... 737-...  
## 2  2013   1998     1     1     533     850 UA      1714 N24211 Fixe... 737-...  
## 3  2013   1990     1     1     542     923 AA      1141 N619AA Fixe... 757-...  
## 4  2013   2012     1     1     544    1004 B6       725 N804JB Fixe... A320-...  
## # i 336,772 more rows
```

What is `year.x`? What is `year.y`?

Summary of key verbs so far

Key verbs

Import

readr

1. `read_csv`
2. `write_csv`

readxl

1. `read_excel`

Tidy

tidyr

1. `pivot_longer`
2. `pivot_wider`
3. `separate_wider_delim`

Join

dplyr

1. `left_join`
2. `right_join`
3. `inner_join`
4. `full_join`
5. `semi_join`
6. `anti_join`

Transform

dplyr

1. `filter`
2. `arrange`
3. `select`
4. `mutate`
5. `summarize`

Logic

Logical vectors

What values can the logical data type take?

Logical values can be **TRUE**, **FALSE**, or **NA**

What are **logical vectors**?

Logical vectors are just vectors ("columns") that only contain **TRUE**, **FALSE**, or **NA**

Logical vectors

Can you think of any logical vectors we have worked with so far?

While we don't often see logical vectors in raw data, we use them all the time!

Example: every time we make comparisons to `filter()` data we create *transient* logical variables that are computed, used, and then thrown away

Transient logical vectors

We create a transient logical vector when we filter `flights` to Miami:

```
library(nycflights13)
flights |>
  select(carrier, flight, dest) |>
  filter(dest == "MIA")
```

```
## # A tibble: 11,728 × 3
##   carrier flight dest
##   <chr>    <int> <chr>
## 1 AA      1141 MIA
## 2 AA      1895 MIA
## 3 UA      1077 MIA
## 4 AA      1837 MIA
## 5 DL      2003 MIA
## 6 AA      2279 MIA
## 7 AA      2267 MIA
## 8 DL      1843 MIA
## 9 AA       443 MIA
## 10 DL     2143 MIA
## # i 11,718 more rows
```

filter(dest == "MIA"): under the hood

```
# create a logical vector from a comparison
flights |>
  select(carrier, flight, dest) |>
  mutate(welcome_to_miami = dest == "MIA")
```

```
## # A tibble: 336,776 × 4
##   carrier flight dest welcome_to_miami
##   <chr>    <int> <chr> <lgl>
## 1 UA      1545 IAH  FALSE
## 2 UA      1714 IAH  FALSE
## 3 AA      1141 MIA  TRUE
## 4 B6       725 BQN  FALSE
## 5 DL       461 ATL  FALSE
## 6 UA      1696 ORD  FALSE
## 7 B6       507 FLL  FALSE
## 8 EV      5708 IAD  FALSE
## 9 B6        79 MCO  FALSE
## 10 AA      301 ORD  FALSE
## # i 336,766 more rows
```

```
flights |>
  select(carrier, flight, dest) |>
  mutate(welcome_to_miami = dest == "MIA") |>
  filter(welcome_to_miami) # then filter
```

```
## # A tibble: 11,728 × 4
##   carrier flight dest welcome_to_miami
##   <chr>    <int> <chr> <lgl>
## 1 AA      1141 MIA  TRUE
## 2 AA      1895 MIA  TRUE
## 3 UA      1077 MIA  TRUE
## 4 AA      1837 MIA  TRUE
## 5 DL      2003 MIA  TRUE
## 6 AA      2279 MIA  TRUE
## 7 AA      2267 MIA  TRUE
## 8 DL      1843 MIA  TRUE
## 9 AA       443 MIA  TRUE
## 10 DL      2143 MIA  TRUE
## # i 11,718 more rows
```

Comparisons

Numeric comparisons like `<`, `<=`, `>`, `>=`, `!=`, and `==` can be used to create logical vectors

As we have seen, `==` and `!=` are useful for comparing characters (i.e., strings)

is.na()

`is.na()` is a useful function for checking whether something is **NA**

Why use `is.na(x)` when we could just use `x == NA`?

```
x <- 2850 + 5850  
is.na(x)
```

```
## [1] FALSE
```

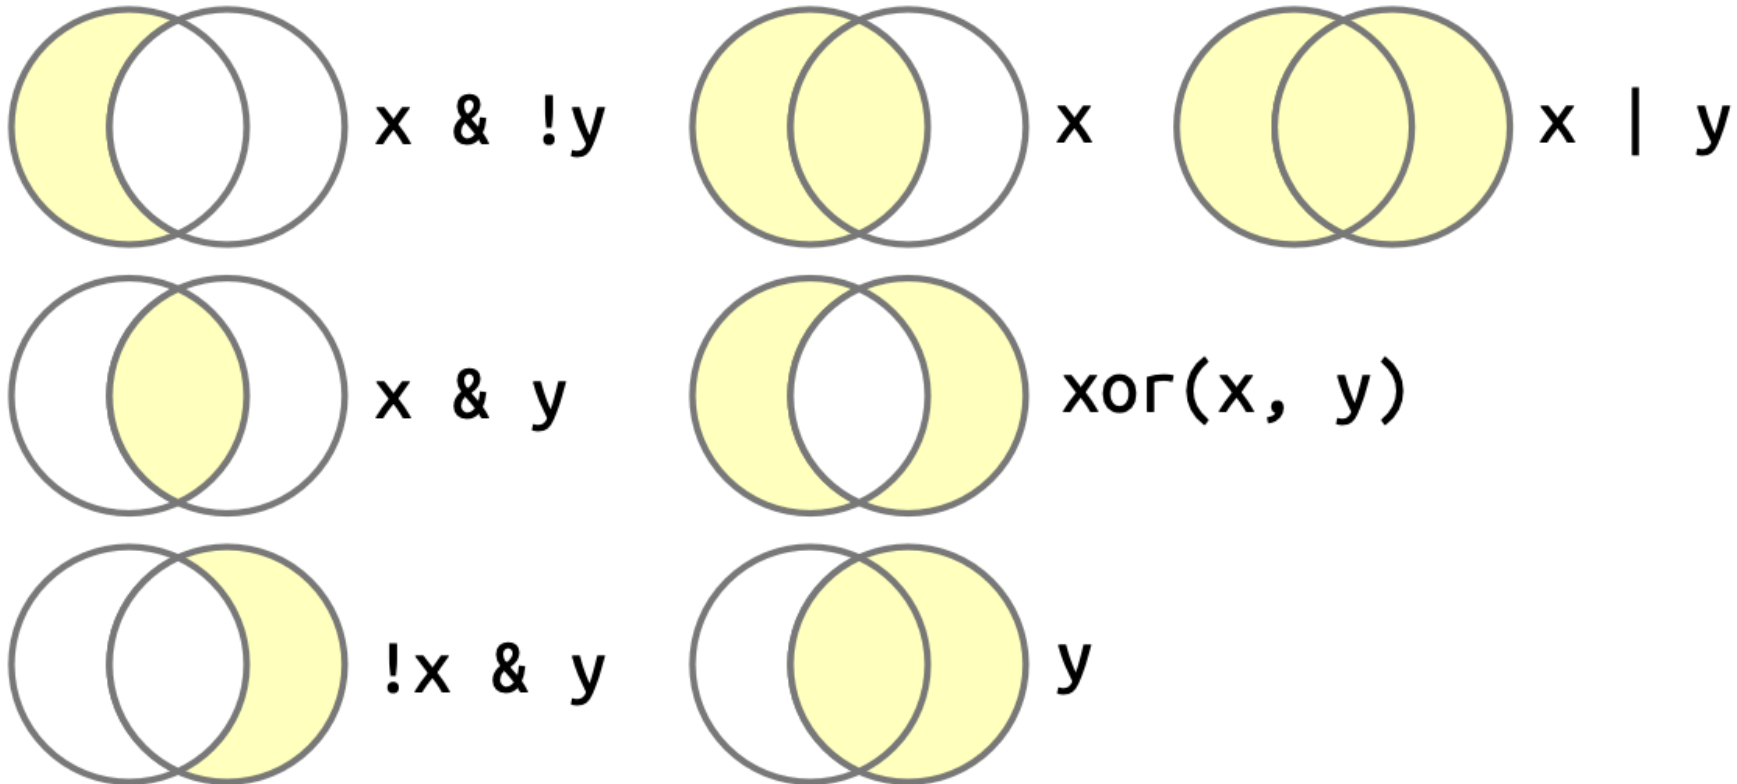
```
x == NA
```

```
## [1] NA
```

May seem odd but it makes sense when you think about concrete comparisons

Boolean algebra

Use Boolean algebra to combine comparisons / logical vectors



Boolean operator examples

```
flights |>
  select(carrier, flight, dest) |>
  filter(dest == "MIA" | dest == "MYR")
```

```
## # A tibble: 11,787 × 3
##   carrier flight dest
##   <chr>    <int> <chr>
## 1 AA        1141 MIA
## 2 AA        1895 MIA
## 3 UA        1077 MIA
## 4 AA        1837 MIA
## 5 DL        2003 MIA
## 6 AA        2279 MIA
## 7 AA        2267 MIA
## 8 DL        1843 MIA
## 9 AA         443 MIA
## 10 EV       4412 MYR
## # i 11,777 more rows
```

```
flights |>
  select(carrier, flight, dest) |>
  filter(carrier == "AA" & dest == "MIA")
```

```
## # A tibble: 7,234 × 3
##   carrier flight dest
##   <chr>    <int> <chr>
## 1 AA        1141 MIA
## 2 AA        1895 MIA
## 3 AA        1837 MIA
## 4 AA        2279 MIA
## 5 AA        2267 MIA
## 6 AA         443 MIA
## 7 AA         647 MIA
## 8 AA        2099 MIA
## 9 AA        1623 MIA
## 10 AA        2253 MIA
## # i 7,224 more rows
```

%in%

`x %in% y` is a useful shortcut for identifying whether a value in `x` is contained in `y`

```
flights |>
  select(carrier, flight, dest) |>
  filter(dest %in% c("MIA", "MYR"))
```

```
## # A tibble: 11,787 × 3
##   carrier flight dest
##   <chr>     <int> <chr>
## 1 AA         1141 MIA
## 2 AA         1895 MIA
## 3 UA         1077 MIA
## 4 AA         1837 MIA
## 5 DL         2003 MIA
## 6 AA         2279 MIA
## 7 AA         2267 MIA
## 8 DL         1843 MIA
## 9 AA          443 MIA
## 10 EV        4412 MYR
## # i 11,777 more rows
```

Numeric operations on logical vectors

Numeric operations treat **TRUE** as **1** and **FALSE** as **0**:

```
x <- c(TRUE, TRUE, FALSE, FALSE, FALSE)
```

```
sum(x)
```

```
## [1] 2
```

```
mean(x)
```

```
## [1] 0.4
```

```
min(x)
```

```
## [1] 0
```

```
max(x)
```

```
## [1] 1
```

This can be handy when doing calculations that depend on conditions

Conditional transformations

`if_else()` can be used to do things based on a binary condition

```
flights |> filter(dest == "MIA") |>
  select(carrier, flight, dest, sched_dep_time) |>
  mutate(too_early = if_else(sched_dep_time < 800, "too early!", "okay"))
```

```
## # A tibble: 11,728 × 5
##   carrier flight dest   sched_dep_time too_early
##   <chr>    <int> <chr>         <int> <chr>
## 1 AA        1141 MIA             540 too early!
## 2 AA        1895 MIA             610 too early!
## 3 UA        1077 MIA             607 too early!
## 4 AA        1837 MIA             610 too early!
## 5 DL        2003 MIA             700 too early!
## 6 AA        2279 MIA             700 too early!
## 7 AA        2267 MIA             755 too early!
## 8 DL        1843 MIA             800 okay
## 9 AA         443 MIA             715 too early!
## 10 DL       2143 MIA             900 okay
## # i 11,718 more rows
```

Conditional transformations

`case_when()` is a more flexible approach that allows many different conditions

```
flights |>
  filter(dest == "MIA") |>
  select(carrier, flight, sched_dep_time) |>
  mutate(too_early = case_when(
    sched_dep_time < 600 ~ "too early!",
    sched_dep_time < 800 ~ "still early",
    sched_dep_time <= 2000 ~ "okay",
    sched_dep_time > 2000 ~ "late"
  ))
```

```
## # A tibble: 11,728 × 4
##   carrier flight sched_dep_time too_early
##   <chr>     <int>         <int> <chr>
## 1 AA         1141             540 too early!
## 2 AA         1895             610 still early
## 3 UA         1077             607 still early
## 4 AA         1837             610 still early
## 5 DL         2003             700 still early
## 6 AA         2279             700 still early
## 7 AA         2267             755 still early
## 8 DL         1843             800 okay
## 9 AA          443             715 still early
## 10 DL        2143             900 okay
## # i 11,718 more rows
```

Conditions are evaluated in order

`condition ~ output` syntax is new; watch out for overlapping conditions!

example-04-2