

Lab-02

Write your name here

2024-01-30

Preface

The goal of this assignment is to help you gain familiarity with data frames – think “spread-sheets” – and how to use **dplyr** functions to transform data. In this lab we are providing some code snippets to serve as “scaffolding” to help guide you through each step. As always, please come to office hours and reach out to your teaching staff if you have any questions.

In some places the scaffolded code provides the argument **FALSE** to functions (e.g., **filter(FALSE)** in question 1). We did this to allow the file to render even before you do any coding. Please replace **FALSE** with your code, just as you have been replacing **-----** with your code in the examples.

We will work with the data table **flights** provided in the package **nycflights13**. The data table includes all domestic flights that departed NYC (i.e. JFK, LGA or EWR) in 2013. It has 19 variables. Details of the package **nycflights13** are available [here](#).

```
head(flights)
```

```
# A tibble: 6 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
1  2013     1     1     517             515           2     830           819
2  2013     1     1     533             529           4     850           830
3  2013     1     1     542             540           2     923           850
4  2013     1     1     544             545          -1    1004          1022
5  2013     1     1     554             600          -6     812           837
6  2013     1     1     554             558          -4     740           728
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dtm>
```

1. In this data set, `arr_delay` is a variable that records the arrival delays in minutes. Negative times represent early arrivals. Use `dplyr::filter` to find: (1) the flights that arrived more than two hours late, and (2) the flights that arrived earlier than scheduled. What is the proportion of flights that arrived more than two hours late? What is the proportion of flights that arrived earlier than scheduled time?

```
# Use dplyr::filter to find and count the flights that arrived more than two hours late
two_hour_late <- flights |>
  filter(FALSE) |> # Replace FALSE with your code
  count()

# Use dplyr::filter to find and count the flights that arrived earlier than scheduled
early_arr <- flights |>
  filter(FALSE) |> # Replace FALSE with your code
  count()

# Count the total number of flights
total <- count(flights)
```

A proportion of 0 of the flights arrived more than two hours late. A proportion of 0 of the flights arrived earlier than scheduled time.

2. How many flights have a missing `dep_time`? What other variables are missing? What might these rows represent?

```
flights |>
  filter(FALSE) |> # Replace FALSE with your code
  count()
```

```
# A tibble: 1 x 1
      n
<int>
1     0
```

These rows probably represent...

3. Use at least two ways to select variables of `dep_time`, `sched_dep_time`, `dep_delay`, `arr_time`, `sched_arr_time`, `arr_delay`. Put `arr_delay` in the first column.

```
# Method 1
# Your code goes here

# Method 2
# Your code goes here
```

4. Use `dplyr::arrange` to sort flights by arrival delays in descending order. How long was the worst arrival delay?

```
worst_delay <- flights |>
  arrange(desc(FALSE)) |> # sort flights by arrival delays in descending order
  filter(FALSE) |> # choose the first row
  pull(arr_delay)
```

The worst arrival delay was inline code minutes.

5. Select `air_time` and `distance`. Generate a new variable `speed` that is calculated as `distance` divided by `air_time` (in miles/min). Then create a variable `mph` that contains speed in miles/hour.

```
flights |>
  select() |> # select `air_time` and `distance` here
  mutate( # create a new variable `speed`
    ) # create a new variable `mph`
```

```
# A tibble: 336,776 x 0
```

6. Select dep_time. Currently dep_time is convenient to look at, but hard to compute with because it is not really a continuous number. Convert it to a represent the number of minutes since midnight. Add a new column dep_time_min to store the converted values.

```
# Your code goes here
flights |>
  select() |>
  mutate(dep_hour = FALSE) |> # split out the hour-digits
  mutate(dep_min = FALSE) |> # split out the minute-digits
  mutate(dep_time_min = dep_hour * 60 + dep_min) # dep_time in min since midnight
```

```
# A tibble: 336,776 x 3
  dep_hour dep_min dep_time_min
  <lgl>    <lgl>         <dbl>
1 FALSE   FALSE           0
2 FALSE   FALSE           0
3 FALSE   FALSE           0
4 FALSE   FALSE           0
5 FALSE   FALSE           0
6 FALSE   FALSE           0
7 FALSE   FALSE           0
8 FALSE   FALSE           0
9 FALSE   FALSE           0
10 FALSE  FALSE           0
# i 336,766 more rows
```

7. Calculate the average arrival delay by carrier. Which carrier has the worst delays?

```
worst_delay_carrier <- flights |>
  group_by() |>
  summarize() |>
  slice_max(FALSE) # Replace FALSE with your code
```

Carrier inline code has the worst delays.

8. What hour(s) of the day should you fly if you want to avoid delays as much as possible?

```
# Your code goes here
```

The best time to fly to avoid delays is...