

Joins and logic

Week 4

AEM 2850 / 5850 : R for Business Analytics

Cornell Dyson

Spring 2023

Acknowledgements: Grant McDermott, Jenny Bryan, R4DS (2e)

Announcements

Reminders:

- Submit assignments via canvas
 - Lab-03 was due yesterday (Monday) at 11:59pm

Questions before we get started?

Plan for this week

Prologue

Joins (Tuesday)

- example-04-1

Logic (Thursday)

- example-04-2

Prologue

What sports do we watch?



Take a guess: what's the most popular spectator sport among classmates?

```
## [1] "football" "basketball" "volleyball" "soccer" NA
## [6] "baseball" "football" "basketball" "basketball" "football"
## [11] "gymnastics" "gymnastics" "squash" NA "basketball"
## [16] "gymnastics" "tennis" "baseball" "basketball" "basketball"
## [21] NA "soccer" "basketball" "baseball" "tennis"
## [26] "polo" "hockey" "basketball" "gymnastics" "soccer"
## [31] "tennis" NA "soccer" "cricket" "tennis"
## [36] "tennis" "badminton" "football" "football"
```

Let's **count** and **arrange** to get the top 3:

```
## # A tibble: 3 × 2
##   sport      n
##   <chr>    <int>
## 1 basketball 8
## 2 football  5
## 3 tennis    5
```

R can be used for sports analytics, too!




Joins

What are relational data?

Multiple tables of data with pairwise relations

Relations across >2 data tables are determined by the relations between each pair

Relational data example: nycflights13



Relational data verbs from dplyr

1. **Mutating joins**: add new variables

- `left_join()`
- `right_join()`
- `inner_join()`
- `full_join()`

2. **Filtering joins**: filter observations

- `semi_join()`
- `anti_join()`

Joins

Let's learn these join commands using two small data frames

superheroes

```
## # A tibble: 7 × 3
##   name      alignment publisher
##   <chr>    <chr>      <chr>
## 1 Magneto  bad         Marvel
## 2 Storm    good        Marvel
## 3 Mystique bad         Marvel
## 4 Batman   good        DC
## 5 Joker    bad         DC
## 6 Catwoman bad         DC
## 7 Hellboy  good        Dark Horse Comics
```

publishers

```
## # A tibble: 3 × 2
##   publisher year_founded
##   <chr>         <int>
## 1 DC             1934
## 2 Marvel         1939
## 3 Image          1992
```

1) dplyr::left_join(x, y)

```
left_join(superheroes, publishers)
```

```
## Joining with `by = join_by(publisher)`
```

```
## # A tibble: 7 × 4
```

	name	alignment	publisher	year_founded
	<chr>	<chr>	<chr>	<int>
## 1	Magneto	bad	Marvel	1939
## 2	Storm	good	Marvel	1939
## 3	Mystique	bad	Marvel	1939
## 4	Batman	good	DC	1934
## 5	Joker	bad	DC	1934
## 6	Catwoman	bad	DC	1934
## 7	Hellboy	good	Dark Horse Comics	NA

`left_join` is a **mutating join**: it adds variables to `x`

`left_join` returns all rows from `x`

2) dplyr::right_join(x, y)

```
right_join(superheroes, publishers)
```

```
## Joining with `by = join_by(publisher)`
```

```
## # A tibble: 7 × 4
```

```
##   name      alignment publisher year_founded
##   <chr>      <chr>      <chr>         <int>
## 1 Magneto   bad         Marvel         1939
## 2 Storm     good        Marvel         1939
## 3 Mystique  bad         Marvel         1939
## 4 Batman    good        DC             1934
## 5 Joker     bad         DC             1934
## 6 Catwoman  bad         DC             1934
## 7 <NA>      <NA>        Image          1992
```

`right_join` is a **mutating join**: it adds variables to `y`

`right_join` returns all rows from `y`

3) dplyr::inner_join(x, y)

```
inner_join(superheroes, publishers)
```

```
## Joining with `by = join_by(publisher)`
```

```
## # A tibble: 6 × 4
```

```
##   name      alignment publisher year_founded
##   <chr>    <chr>      <chr>      <int>
## 1 Magneto  bad         Marvel      1939
## 2 Storm    good        Marvel      1939
## 3 Mystique bad         Marvel      1939
## 4 Batman   good         DC           1934
## 5 Joker    bad          DC           1934
## 6 Catwoman bad          DC           1934
```

How is `inner_join` different from `left_join` and `right_join`?

`inner_join` returns all rows in `x` **AND** `y`

4) dplyr::full_join(x, y)

```
full_join(superheroes, publishers) # how many rows do you think this will produce?
```

```
## Joining with `by = join_by(publisher)`
```

```
## # A tibble: 8 × 4
```

	name	alignment	publisher	year_founded
	<chr>	<chr>	<chr>	<int>
## 1	Magneto	bad	Marvel	1939
## 2	Storm	good	Marvel	1939
## 3	Mystique	bad	Marvel	1939
## 4	Batman	good	DC	1934
## 5	Joker	bad	DC	1934
## 6	Catwoman	bad	DC	1934
## 7	Hellboy	good	Dark Horse Comics	NA
## 8	<NA>	<NA>	Image	1992

full_join returns all rows in x **OR** y

5) dplyr::semi_join(x, y)

```
superheroes
```

```
## # A tibble: 7 × 3
##   name      alignment publisher
##   <chr>    <chr>    <chr>
## 1 Magneto  bad      Marvel
## 2 Storm    good     Marvel
## 3 Mystique bad      Marvel
## 4 Batman   good     DC
## 5 Joker    bad      DC
## 6 Catwoman bad      DC
## 7 Hellboy  good     Dark Horse Comics
```

```
semi_join(superheroes, publishers)
```

```
## Joining with `by = join_by(publisher)`

## # A tibble: 6 × 3
##   name      alignment publisher
##   <chr>    <chr>    <chr>
## 1 Magneto  bad      Marvel
## 2 Storm    good     Marvel
## 3 Mystique bad      Marvel
## 4 Batman   good     DC
## 5 Joker    bad      DC
## 6 Catwoman bad      DC
```

`semi_join` is a **filtering join**: it keeps observations in `x` that have a match in `y`

Note that the variables do not change

6) dplyr::anti_join(x, y)

```
superheroes
```

```
## # A tibble: 7 × 3
##   name      alignment publisher
##   <chr>    <chr>      <chr>
## 1 Magneto  bad         Marvel
## 2 Storm    good        Marvel
## 3 Mystique bad         Marvel
## 4 Batman   good        DC
## 5 Joker    bad         DC
## 6 Catwoman bad         DC
## 7 Hellboy  good        Dark Horse Comics
```

```
anti_join(superheroes, publishers)
```

```
## Joining with `by = join_by(publisher)`

## # A tibble: 1 × 3
##   name      alignment publisher
##   <chr>    <chr>      <chr>
## 1 Hellboy  good        Dark Horse Comics
```

`anti_join` is a **filtering join**: it keeps obs. in `x` that **DO NOT** have a match in `y`

Note that the variables do not change

Key variables

How do `dplyr` join commands know what variables to use as **keys**?

By default, `*_join()` uses all variables that are common across `x` and `y`

```
intersect(names(superheroes), names(publishers)) # variable used for matching before
```

```
## [1] "publisher"
```

Or, we can specify what to join by: `*_join(..., by = join_by(publisher))`

Note: before `dplyr` 1.1.0, the syntax was: `*_join(..., by = "publisher")`

Exploring keys

```
library(nycflights13) # let's explore keys using the nycflights13 data
flights |> print(n = 8) # print the first 8 rows of flights
```

```
## # A tibble: 336,776 × 19
##   year month   day dep_time sched_dep...1 dep_d...2 arr_t...3 sched...4 arr_d...5 carrier
##   <int> <int> <int>   <int>         <int>    <dbl>    <int>    <int>    <dbl> <chr>
## 1  2013     1     1     517           515        2      830      819      11 UA
## 2  2013     1     1     533           529        4      850      830     20 UA
## 3  2013     1     1     542           540        2      923      850     33 AA
## 4  2013     1     1     544           545       -1     1004     1022    -18 B6
## 5  2013     1     1     554           600       -6      812      837    -25 DL
## 6  2013     1     1     554           558       -4      740      728     12 UA
## 7  2013     1     1     555           600       -5      913      854     19 B6
## 8  2013     1     1     557           600       -3      709      723    -14 EV
## # ... with 336,768 more rows, 9 more variables: flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>, and abbreviated variable names
## #   1sched_dep_time, 2dep_delay, 3arr_time, 4sched_arr_time, 5arr_delay
```

Exploring keys

```
planes # print the first 10 rows of planes
```

```
## # A tibble: 3,322 × 9
##   tailnum year type      manuf...1 model engines seats speed engine
##   <chr>   <int> <chr>      <chr>   <chr>   <int> <int> <int> <chr>
## 1 N10156  2004 Fixed wing multi engi... EMBRAER EMB-...     2    55    NA Turbo...
## 2 N102UW  1998 Fixed wing multi engi... AIRBUS... A320...     2   182    NA Turbo...
## 3 N103US  1999 Fixed wing multi engi... AIRBUS... A320...     2   182    NA Turbo...
## 4 N104UW  1999 Fixed wing multi engi... AIRBUS... A320...     2   182    NA Turbo...
## 5 N10575  2002 Fixed wing multi engi... EMBRAER EMB-...     2    55    NA Turbo...
## 6 N105UW  1999 Fixed wing multi engi... AIRBUS... A320...     2   182    NA Turbo...
## 7 N107US  1999 Fixed wing multi engi... AIRBUS... A320...     2   182    NA Turbo...
## 8 N108UW  1999 Fixed wing multi engi... AIRBUS... A320...     2   182    NA Turbo...
## 9 N109UW  1999 Fixed wing multi engi... AIRBUS... A320...     2   182    NA Turbo...
## 10 N110UW 1999 Fixed wing multi engi... AIRBUS... A320...     2   182    NA Turbo...
## # ... with 3,312 more rows, and abbreviated variable name 1manufacturer
```

Let's perform a left join on flights and planes

```
left_join(flights, planes) |>  
  select(year:dep_time, arr_time, carrier:tailnum, type, model) |> # keep text to one slide  
  print(n = 5) # just to save vertical space on the slide
```

```
## Joining with `by = join_by(year, tailnum)`
```

```
## # A tibble: 336,776 × 10
```

```
##   year month   day dep_time arr_time carrier flight tailnum type  model  
##   <int> <int> <int>   <int>   <int>   <chr>   <int> <chr>   <chr> <chr>  
## 1  2013     1     1     517     830   UA      1545 N14228 <NA> <NA>  
## 2  2013     1     1     533     850   UA      1714 N24211 <NA> <NA>  
## 3  2013     1     1     542     923   AA      1141 N619AA <NA> <NA>  
## 4  2013     1     1     544    1004   B6       725 N804JB <NA> <NA>  
## 5  2013     1     1     554     812   DL       461 N668DN <NA> <NA>
```

```
## # ... with 336,771 more rows
```

Uh-oh! What's up with **type** and **model**?

Uh-oh!

As before, `dplyr` guessed which columns to join on

It uses columns with the same name:

```
## Joining, by = c("year", "tailnum")
```


Does anyone see a potential problem here?

The variable `year` does not have a consistent meaning across the datasets

In `flights` it refers to the *year of flight*, in `planes` it refers to *year of construction*

Luckily we can avoid this by using the argument `by = join_by(...)`

What should we join flights and planes by?



Specifying join keys

We just need to be explicit in the join call by using the **by** argument

```
left_join(flights,  
          planes |> rename(year_built = year), # not necessary w/ below line, but helpful  
          by = join_by(tailnum) # be specific about the joining column  
        ) |>  
select(year, month:dep_time, carrier, flight, tailnum, year_built, type, model) |>  
print(n = 5) # just to save vertical space on the slide
```

```
## # A tibble: 336,776 × 10  
##   year month   day dep_time carrier flight tailnum year_built type      model  
##   <int> <int> <int>   <int> <chr>   <int> <chr>      <int> <chr>    <chr>  
## 1  2013     1     1     517 UA      1545 N14228     1999 Fixed wing... 737-...  
## 2  2013     1     1     533 UA      1714 N24211     1998 Fixed wing... 737-...  
## 3  2013     1     1     542 AA      1141 N619AA     1990 Fixed wing... 757-...  
## 4  2013     1     1     544 B6       725 N804JB     2012 Fixed wing... A320...  
## 5  2013     1     1     554 DL       461 N668DN     1991 Fixed wing... 757-...  
## # ... with 336,771 more rows
```


Specifying join keys

What happens if we don't rename `year` before this join?

```
left_join(flights,  
  planes, # not renaming "year" to "year_built" this time  
  by = join_by(tailnum)  
) |>  
select(contains("year"), month:dep_time, arr_time, carrier, flight, tailnum, type, model) |>  
print(n = 4) # just to save vertical space on the slide
```

```
## # A tibble: 336,776 × 11  
##   year.x year.y month   day dep_time arr_time carrier flight tailnum type  model  
##   <int> <int> <int> <int>   <int>   <int> <chr>   <int> <chr>  <chr> <chr>  
## 1  2013   1999     1     1     517     830 UA      1545 N14228 Fixe... 737-...  
## 2  2013   1998     1     1     533     850 UA      1714 N24211 Fixe... 737-...  
## 3  2013   1990     1     1     542     923 AA      1141 N619AA Fixe... 757-...  
## 4  2013   2012     1     1     544    1004 B6       725 N804JB Fixe... A320-...  
## # ... with 336,772 more rows
```

What is `year.x`? What is `year.y`?

Summary of key verbs so far

Key verbs

Import

readr

1. `read_csv`
2. `write_csv`

readxl

1. `read_excel`

Tidy

tidyr

1. `pivot_longer`
2. `pivot_wider`
3. `separate`
4. `unite`

Join

dplyr

1. `left_join`
2. `right_join`
3. `inner_join`
4. `full_join`
5. `semi_join`
6. `anti_join`

Transform

dplyr

1. `filter`
2. `arrange`
3. `select`
4. `mutate`
5. `summarize`

link to example-04-1

Logic

Logic slides will be added for Thursday

See <https://aem2850.toddgerarden.com/content/04-content/> for the updated version of these slides

link to example-04-2