

Strings and regular expressions

Week 5

AEM 2850 / 5850 : R for Business Analytics

Cornell Dyson

Spring 2024

Acknowledgements: R4DS (2e)

Announcements

No lab due next Monday (February break)

Prelim 1 will be next Thursday, February 29 in class

- Will cover everything so far **including this week's material**
- We will provide more guidance Thursday and via canvas
- Please contact me and SDS **as soon as possible** if you are eligible for testing accommodations and are not yet registered

Questions before we get started?

Plan for today

Prologue: clarifying group operations

Working with strings in R

Regular expressions

Working with regular expressions in R

Prologue

Group operations

Lab-04 included the following question:

3. Calculate the average housing price for each of the cities that are in `tidy_txhouse` and in `tidy_txpop` over the period 2010 through 2012. Which city has the highest average housing price? Which city has the lowest average housing price?

This is a good example question to provide more detail on `group_by()` and `ungroup()`, which I mentioned in passing during week 2

Reminder: dplyr::group_by

`summarize` is particularly useful in combination with `group_by`:

```
semi_join(tidy_txhouse, tx_pop_city) |> # for cities in tidy_txhouse and tx_pop_city
  filter(2010 <= year & year <= 2012) |> # for years 2010-2012
  group_by(city) |> # for each city
  summarize(mean_price = mean(price)) # compute the mean price
```

```
## # A tibble: 34 × 2
##   city          mean_price
##   <chr>          <dbl>
## 1 Abilene        113614.
## 2 Amarillo       128192.
## 3 Arlington     128058.
## 4 Austin        193811.
## 5 Beaumont      128431.
## 6 Brownsville    96569.
## 7 Corpus Christi 137883.
## 8 Dallas         161706.
## 9 El Paso        135117.
## 10 Fort Worth    114217.
## # i 24 more rows
```

Groups are persistent (sort of)

```
semi_join(tidy_txhouse, tx_pop_city) |>
  filter(2010 <= year & year <= 2012) |>
  group_by(city) |>
  print(n = 5)
```

```
## # A tibble: 102 × 4
```

```
## # Groups:   city [34]
```

```
##   city      year sales  price
##   <chr>    <dbl> <dbl>   <dbl>
```

```
## 1 Abilene   2010   1590 112000
```

```
## 2 Abilene   2011   1719 114492.
```

```
## 3 Abilene   2012   2016 114350
```

```
## 4 Amarillo  2010   2556 125117.
```

```
## 5 Amarillo  2011   2705 126658.
```

```
## # i 97 more rows
```

```
semi_join(tidy_txhouse, tx_pop_city) |>
  filter(2010 <= year & year <= 2012) |>
  group_by(city) |>
  summarize(mean_price = mean(price)) |>
  print(n = 5)
```

```
## # A tibble: 34 × 2
```

```
##   city      mean_price
```

```
##   <chr>         <dbl>
```

```
## 1 Abilene      113614.
```

```
## 2 Amarillo     128192.
```

```
## 3 Arlington    128058.
```

```
## 4 Austin       193811.
```

```
## 5 Beaumont     128431.
```

```
## # i 29 more rows
```

No groups?!

Groups are persistent (sort of)

`summarize()` drops the last group variable, so the output is not grouped by `city`

By contrast `mutate()` does not unroll group variables unless you explicitly ask it to

```
semi_join(tidy_txhouse, tx_pop_city) |>
  filter(2010 <= year & year <= 2012) |>
  group_by(city) |>
  mutate(mean_price = mean(price))
```

```
## # A tibble: 102 × 5
## # Groups:   city [34]
##   city      year sales  price mean_price
##   <chr>    <dbl> <dbl>   <dbl>      <dbl>
## 1 Abilene   2010   1590 112000    113614.
## 2 Abilene   2011   1719 114492.    113614.
## 3 Abilene   2012   2016 114350    113614.
## 4 Amarillo  2010   2556 125117.    128192.
## 5 Amarillo  2011   2705 126658.    128192.
## 6 Amarillo  2012   2933 132800     128192.
## 7 Arlington 2010   3883 129717.    128058.
## 8 Arlington 2011   3719 124267.    128058.
## 9 Arlington 2012   4248 130192.    128058.
## 10 Austin   2010  19872 189658.    193811.
## # i 92 more rows
```


Groups affect filter operations!

```
semi_join(tidy_txhouse, tx_pop_city) |>
  filter(2010 <= year & year <= 2012) |>
  group_by(city) |>
  summarize(mean_price = mean(price)) |>
  filter(mean_price == max(mean_price))
```

```
## # A tibble: 1 × 2
##   city    mean_price
##   <chr>      <dbl>
## 1 Austin    193811.
```

Looks good!

```
semi_join(tidy_txhouse, tx_pop_city) |>
  filter(2010 <= year & year <= 2012) |>
  group_by(city) |>
  mutate(mean_price = mean(price)) |>
  filter(mean_price == max(mean_price)) |>
  head(5)
```

```
## # A tibble: 5 × 5
## # Groups:   city [2]
##   city    year sales    price mean_price
##   <chr>  <dbl> <dbl>    <dbl>      <dbl>
## 1 Abilene  2010   1590  112000    113614.
## 2 Abilene  2011   1719  114492.    113614.
## 3 Abilene  2012   2016  114350    113614.
## 4 Amarillo 2010   2556  125117.    128192.
## 5 Amarillo 2011   2705  126658.    128192.
```

What went wrong?

Groups affect filter operations!

One way to get around this is to use `ungroup()` or `group_by()` with no arguments

Now we get Austin and only Austin

But we still get multiple rows, when we only need/want one

```
semi_join(tidy_txhouse, tx_pop_city) |>
  filter(2010 <= year & year <= 2012) |>
  group_by(city) |>
  mutate(mean_price = mean(price)) |>
  ungroup() |>
  filter(mean_price == max(mean_price)) |>
  head(5)
```

```
## # A tibble: 3 × 5
##   city    year sales  price mean_price
##   <chr>  <dbl> <dbl>   <dbl>      <dbl>
## 1 Austin  2010  19872 189658.   193811.
## 2 Austin  2011  21208 190033.   193811.
## 3 Austin  2012  25521 201742.   193811.
```

Working with strings in R

Strings are nothing new

```
flights |>
  select(carrier, tailnum, origin, dest)
```

```
## # A tibble: 336,776 × 4
##   carrier tailnum origin dest
##   <chr>   <chr>   <chr> <chr>
## 1 UA      N14228   EWR   IAH
## 2 UA      N24211   LGA   IAH
## 3 AA      N619AA   JFK   MIA
## 4 B6      N804JB   JFK   BQN
## 5 DL      N668DN   LGA   ATL
## 6 UA      N39463   EWR   ORD
## 7 B6      N516JB   EWR   FLL
## 8 EV      N829AS   LGA   IAD
## 9 B6      N593JB   JFK   MCO
## 10 AA     N3ALAA   LGA   ORD
## # i 336,766 more rows
```

```
tx_pop_city |>
  select(city)
```

```
## # A tibble: 1,217 × 1
##   city
##   <chr>
## 1 Abbott
## 2 Abernathy
## 3 Abilene
## 4 Ackerly
## 5 Addison
## 6 Adrian
## 7 Agua Dulce
## 8 Alamo
## 9 Alamo Heights
## 10 Alba
## # i 1,207 more rows
```

Strings in R

Strings are also referred to as "characters" (abbreviated `chr`)

Strings can be stored in many ways:

- Vectors
- **Data frame columns**
- Elements in a list

So far we have used them as we would any other data

But now we'll learn to filter on, modify, or analyze "functions" of strings

The stringr package

`stringr` is loaded as part of the core tidyverse

All `stringr` functions have intuitive names that start with `str_`

We will cover a bunch of handy functions this week:

1. `str_length`
2. `str_to_upper` and `str_to_lower`
3. `str_c` and `str_glue`
4. `str_detect`
5. `str_count`
6. `str_replace`

See `vignette("stringr")` for more

We'll use data from The Office

The **schrute** package contains transcripts of all episodes of **The Office** (US)

```
library(schrute)
theoffice # this data frame is an object from the schrute package
```

```
## # A tibble: 55,130 × 12
##   index season episode episode_name director      writer      character text
##   <int>  <int>   <int>   <chr>      <chr>      <chr>      <chr>   <chr>
## 1      1      1       1     Pilot      Ken Kwapis Ricky Gervais;S... Michael All ...
## 2      2      1       1     Pilot      Ken Kwapis Ricky Gervais;S... Jim      Oh, ...
## 3      3      1       1     Pilot      Ken Kwapis Ricky Gervais;S... Michael So y...
## 4      4      1       1     Pilot      Ken Kwapis Ricky Gervais;S... Jim      Actu...
## 5      5      1       1     Pilot      Ken Kwapis Ricky Gervais;S... Michael All ...
## 6      6      1       1     Pilot      Ken Kwapis Ricky Gervais;S... Michael Yes,...
## 7      7      1       1     Pilot      Ken Kwapis Ricky Gervais;S... Michael I've...
## 8      8      1       1     Pilot      Ken Kwapis Ricky Gervais;S... Pam      Well...
## 9      9      1       1     Pilot      Ken Kwapis Ricky Gervais;S... Michael If y...
## 10     10      1       1     Pilot      Ken Kwapis Ricky Gervais;S... Pam      What?
## # i 55,120 more rows
## # i 4 more variables: text_w_direction <chr>, imdb_rating <dbl>,
## #   total_votes <int>, air_date <chr>
```

1) str_length()

`str_length` tells you the number of characters in a string

```
str_length("supercalifragilisticexpialidocious")
```

```
## [1] 34
```

```
theoffice |>  
  distinct(character) |>  
  slice_head(n = 5) |>  
  mutate(name_length = str_length(character))
```

```
## # A tibble: 5 × 2  
##   character name_length  
##   <chr>         <int>  
## 1 Michael           7  
## 2 Jim                3  
## 3 Pam                3  
## 4 Dwight            6  
## 5 Jan                3
```


2) `str_to_lower()` and `str_to_upper()`

`str_to_lower` converts to lower case

```
str_to_lower("I went to Cornell, you ever heard of it?")
```

```
## [1] "i went to cornell, you ever heard of it?"
```

`str_to_upper` converts to upper case

```
str_to_upper("I went to Cornell, you ever heard of it?")
```

```
## [1] "I WENT TO CORNELL, YOU EVER HEARD OF IT?"
```

Similar functionality for `str_to_title()` and `str_to_sentence()`

These functions are locale dependent (e.g., "en_GB" vs "en_US")

3) str_c()

We have seen `c()` combine arguments into a vector or list

Similarly, `str_c()` combines arguments into a character vector:

```
str_c("a", "b", "c", "1", "2", "3")
```

```
## [1] "abc123"
```

Here, it combined six letters and numbers into a single string

But we can also use it within data frames to combine strings at scale

3) str_c()

Here's an example of multiple columns in a data frame being combined into one:

```
theoffice |> slice_head(n = 1) |>  
  select(character, text)
```

```
## # A tibble: 1 × 2  
##   character text  
##   <chr>      <chr>  
## 1 Michael   All right Jim. Your quarterlies look very good. How are things at t...
```

```
theoffice |> slice_head(n = 1) |>  
  transmute(line = str_c(character, " said: ", text)) # mutate and keep only `line`
```

```
## # A tibble: 1 × 1  
##   line  
##   <chr>  
## 1 Michael said: All right Jim. Your quarterlies look very good. How are things ...
```

3) str_c()

`str_c()` will automatically recycle fixed arguments like "**said:** " that are shorter than `character` and `text`:

```
theoffice |> slice_head(n = 3) |>  
  transmute(line = str_c(character, " said: ", text))
```

```
## # A tibble: 3 × 1  
##   line  
##   <chr>  
## 1 Michael said: All right Jim. Your quarterlies look very good. How are things ...  
## 2 Jim said: Oh, I told you. I couldn't close it. So...  
## 3 Michael said: So you've come to the master for guidance? Is this what you're ...
```

`str_c()` and `str_glue()` work well with `mutate()` and `transmute()` because their output is the same length as their inputs

3) str_glue()

`str_glue()` provides similar functionality, but different syntax:

```
theoffice |> slice_head(n = 3) |>  
  transmute(line = str_glue("{character} said: {text}")) # note the different syntax
```

```
## # A tibble: 3 × 1  
##   line  
##   <glue>  
## 1 Michael said: All right Jim. Your quarterlies look very good. How are things ...  
## 2 Jim said: Oh, I told you. I couldn't close it. So...  
## 3 Michael said: So you've come to the master for guidance? Is this what you're ...
```

Items inside `{ }` are evaluated as if they are outside the quotes

This can be handy when combining many fixed and variable strings

Regular expressions

Regular expressions

What are regular expressions?

A concise, powerful way for describing patterns within strings

Regular expressions are a generic tool, not something specific to R

Let's use the names of some characters from The Office as examples:

```
names <- theoffice |> distinct(character) |> slice_head(n = 10) |> pull(character)
names
```

```
## [1] "Michael"      "Jim"           "Pam"           "Dwight"        "Jan"
## [6] "Michel"       "Todd Packer"  "Phyllis"       "Stanley"       "Oscar"
```

Pattern basics

The simplest patterns consist of literal characters

```
names
```

```
## [1] "Michael"      "Jim"          "Pam"
## [4] "Dwight"       "Jan"          "Michel"
## [7] "Todd Packer"  "Phyllis"      "Stanley"
## [10] "Oscar"
```

`str_view()` is a handy utility to see how patterns match

What do you think this will return?

```
str_view(names, pattern = "J")
```

```
## [2] | <J>im
## [5] | <J>an
```

Literal pattern matches are case-sensitive by default

Meta-characters

Punctuation characters like `.`, `+`, `*`, `[`, `]`, and `?` are **meta-characters** with special meanings

The most common one is `.`, which will match any character

What do you think these statements will return?

```
str_view(names, pattern = "J.m")
```

```
## [2] | <Jim>
```

```
str_view(names, pattern = "J.n")
```

```
## [5] | <Jan>
```

Meta-characters

What do you think these statements will return?

```
str_view(names, pattern = "J..")
```

```
str_view(names, pattern = "J...")
```

```
## [2] | <Jim>  
## [5] | <Jan>
```

Quantifiers

Quantifiers control how many times a pattern can match:

- **?** makes a pattern optional -- it matches 0 or 1 times
- **+** lets a pattern repeat -- it matches at least once
- ***** lets a pattern be optional or repeat

What do you think this statement will return?

```
str_view(names, "M.*l") # match strings with M, then any number of any characters, then l
```

```
## [1] | <Michael>  
## [6] | <Michel>
```

Character classes

`[]` lets you match a set of characters

```
str_view(names, "[aeiou]") # vowels
```

```
## [1] M<i>ch<a><e>l
## [2] J<i>m
## [3] P<a>m
## [4] Dw<i>ght
## [5] J<a>n
## [6] M<i>ch<e>l
## [7] T<o>dd P<a>ck<e>r
## [8] Phyll<i>s
## [9] St<a>nl<e>y
## [10] Osc<a>r
```

`^` inverts character class matches

```
str_view(names, "[^aeiou]") # NOT vowels
```

```
## [1] <M>i<c><h>ae<l>
## [2] <J>i<m>
## [3] <P>a<m>
## [4] <D><w>i<g><h><t>
## [5] <J>a<n>
## [6] <M>i<c><h>e<l>
## [7] <T>o<d><d>< ><P>a<c><k>e<r>
## [8] <P><h><y><l><l>i<s>
## [9] <S><t>a<n><l>e<y>
## [10] <O><s><c>a<r>
```

Alternation

Last one! Hang in there!

Alternation, `|`, allows you to search for one or more alternative patterns

This should seem familiar...

What do you think these statements will return?

```
str_view(names, "J.m|P.m")
```

```
## [2] | <Jim>  
## [3] | <Pam>
```

```
str_view(names, "S.*|O.*")
```

```
## [9] | <Stanley>  
## [10] | <Oscar>
```

More patterns

See [Chapter 15 of R4DS \(2e\)](#) for more on:

- **escaping**: matching meta-characters as if they were literal strings
- **anchors**: match the start or end of a string
- **character classes** (continued)
- **quantifiers** (continued)
- **operator precedence**: parentheses, etc.
- **grouping**: back references, etc.

Working with regular expressions in R

4) str_detect()

`str_detect` can be used to match patterns and return a logical vector

```
first_4_characters <- theoffice |>
  distinct(character) |>
  slice_head(n = 4) |>
  pull(character)
first_4_characters
```

```
## [1] "Michael" "Jim"      "Pam"      "Dwight"
```

```
str_detect(first_4_characters, "Dwight")
```

```
## [1] FALSE FALSE FALSE  TRUE
```

```
str_detect(first_4_characters, "a")
```

```
## [1]  TRUE FALSE  TRUE FALSE
```

How could we fit this into our current workflow?

4) str_detect()

`str_detect` is a powerful way to `filter` a data frame

```
theoffice |> select(season, episode, character, text) |>  
  filter(str_detect(text,      # where to match a pattern  
                  "sale")) # what pattern to match
```

```
## # A tibble: 370 × 4  
##   season episode character text  
##   <int>   <int> <chr>    <chr>  
## 1       1       2 Jim      This is my biggest sale of the year. They love me o...  
## 2       1       2 Jim      Mr. Decker, we didn't lose your sale today, did we?...  
## 3       1       3 Jim      That is a great offer. Thank you. I really think I ...  
## 4       1       3 Jan      From sales?  
## 5       1       4 Michael Look, look, look. I talked to corporate, about prot..  
## 6       1       5 Michael All right, time, time out. Come on, sales, over her..  
## 7       1       6 Jan      Alan and I have created an incentive program to inc..  
## 8       1       6 Jan      We've created an incentive program to increase sale..  
## 9       1       6 Jim      Plus you have so much more to talk to this girl abo..  
## 10      1       6 Stanley I thought that was the incentive prize for the top ...  
## # i 360 more rows
```

4) str_detect()

Literal pattern matches with `str_detect` are case-sensitive

```
theoffice |> select(season, episode, character, text) |>
  filter(str_detect(text,
    "Sale")) # sale and Sale produce different output
```

```
## # A tibble: 28 × 4
##   season episode character      text
##   <int>   <int> <chr>      <chr>
## 1       2     11 Michael    No, no. Salesmen and profit centers.
## 2       2     14 Michael    Old fashioned raid. Sales on Accounting. Y...
## 3       2     14 Michael and Dwight Ahhhh! Whoo hoo! Come on, come on, come on...
## 4       2     14 Michael    Oh, and I'm not? Why would you say that? B...
## 5       2     17 Jim        Dwight was the top salesman of the year at...
## 6       2     17 Michael    Speaker at the Sales Convention. Been ther...
## 7       2     17 Dwight    Saleswoman has a v*gl*n*.
## 8       2     17 Speaker    Next, I'd like to introduce the Dunder Mif...
## 9       2     17 Dwight    Salesman of Northeastern Pennsylvania, I a...
## 10      3      5 Angela    Sales take a long time.
## # i 18 more rows
```

4) str_detect()

You could use multiple calls to `str_detect`, or use alternation:

```
theoffice |> select(season, episode, character, text) |>
  filter(str_detect(text,
    "sale|Sale")) # look for sale OR Sale
```

```
## # A tibble: 392 × 4
##   season episode character text
##   <int>   <int> <chr>    <chr>
## 1       1       2 Jim      This is my biggest sale of the year. They love me o...
## 2       1       2 Jim      Mr. Decker, we didn't lose your sale today, did we?...
## 3       1       3 Jim      That is a great offer. Thank you. I really think I ...
## 4       1       3 Jan      From sales?
## 5       1       4 Michael Look, look, look. I talked to corporate, about prot...
## 6       1       5 Michael All right, time, time out. Come on, sales, over her...
## 7       1       6 Jan      Alan and I have created an incentive program to inc...
## 8       1       6 Jan      We've created an incentive program to increase sale...
## 9       1       6 Jim      Plus you have so much more to talk to this girl abo...
## 10      1       6 Stanley I thought that was the incentive prize for the top ...
## # i 382 more rows
```

4) str_detect()

You could consolidate this: regex parentheses are like in math

```
theoffice |> select(season, episode, character, text) |>
  filter(str_detect(text,
    "(s|S)ale")) # look for sale OR Sale
```

```
## # A tibble: 392 × 4
##   season episode character text
##   <int>   <int> <chr>    <chr>
## 1       1       2 Jim      This is my biggest sale of the year. They love me o...
## 2       1       2 Jim      Mr. Decker, we didn't lose your sale today, did we?...
## 3       1       3 Jim      That is a great offer. Thank you. I really think I ...
## 4       1       3 Jan      From sales?
## 5       1       4 Michael Look, look, look. I talked to corporate, about prot...
## 6       1       5 Michael All right, time, time out. Come on, sales, over her...
## 7       1       6 Jan      Alan and I have created an incentive program to inc...
## 8       1       6 Jan      We've created an incentive program to increase sale...
## 9       1       6 Jim      Plus you have so much more to talk to this girl abo...
## 10      1       6 Stanley I thought that was the incentive prize for the top ...
## # i 382 more rows
```

4) str_detect()

Or use `regex()` to ignore all cases and control other pattern matching details

```
theoffice |> select(season, episode, character, text) |>
  filter(str_detect(text,
    regex("Sale", ignore_case = TRUE))) # ignore case
```

```
## # A tibble: 393 × 4
##   season episode character text
##   <int>   <int> <chr>    <chr>
## 1         1       2 Jim      This is my biggest sale of the year. They love me o...
## 2         1       2 Jim      Mr. Decker, we didn't lose your sale today, did we?...
## 3         1       3 Jim      That is a great offer. Thank you. I really think I ...
## 4         1       3 Jan      From sales?
## 5         1       4 Michael Look, look, look. I talked to corporate, about prot...
## 6         1       5 Michael All right, time, time out. Come on, sales, over her...
## 7         1       6 Jan      Alan and I have created an incentive program to inc...
## 8         1       6 Jan      We've created an incentive program to increase sale...
## 9         1       6 Jim      Plus you have so much more to talk to this girl abo...
## 10        1       6 Stanley I thought that was the incentive prize for the top ...
## # i 383 more rows
```

4) str_detect()

When I say ignore all cases, I mean IGNORE ALL CASES!

```
theoffice |> select(season, episode, character, text) |>
  filter(str_detect(text,
                    regex("sale", ignore_case = TRUE))) |>
  filter(!str_detect(text, "(s|S)ale")) # find non-standard form(s)
```

```
## # A tibble: 1 × 4
##   season episode character text
##   <int>   <int> <chr>      <chr>
## 1       3       3 Dwight    I HAVE EXCELLENT SALES NUMBERS!
```

4) str_detect()

`str_detect` can be combined with familiar functions to summarize data

```
theoffice |>
  filter(str_detect(text, regex("Sale", ignore_case = TRUE))) |>
  count(character, sort = TRUE)
```

```
## # A tibble: 46 × 2
##   character      n
##   <chr>      <int>
## 1 Michael      91
## 2 Dwight       81
## 3 Jim          51
## 4 Andy         31
## 5 Pam          26
## 6 Ryan         10
## 7 Clark         8
## 8 Gabe          7
## 9 David         6
## 10 Angela       5
## # i 36 more rows
```

4) str_detect()

`str_detect` can be combined with familiar functions to summarize data

```
theoffice |>
  filter(str_detect(text,
    regex("that's what she said", ignore_case = TRUE))) |>
  count(character, sort = TRUE)
```

```
## # A tibble: 8 × 2
##   character      n
##   <chr>      <int>
## 1 Michael      23
## 2 Dwight        3
## 3 Jim           2
## 4 Creed         1
## 5 David         1
## 6 Holly         1
## 7 Jan           1
## 8 Pam           1
```


4) str_detect()

`str_detect` with regular expressions can be very powerful

```
theoffice |> select(character, text) |>  
  filter(str_detect(text, "assistant.*manager")) |>  
  slice_head(n = 10)
```

```
## # A tibble: 10 × 2  
##   character text  
##   <chr>      <chr>  
## 1 Dwight    I, but if there were, I'd be protected as assistant regional manag...  
## 2 Dwight    And that's why you have an assistant regional manager.  
## 3 Michael   No, I am the team manager. You can be assistant to the team manage...  
## 4 Dwight    Hey, Pam, I'm assistant regional manager, and I can take care of h...  
## 5 Michael   All right. Well then, you are now acting manager of Dunder Mifflin...  
## 6 Dwight    Uh,... my first sale, my promotion to assistant regional manager, ...  
## 7 Jim       Oh, that's because at first it was a made up position for Dwight, ...  
## 8 Charles   So you're the assistant to the regional manager?  
## 9 Darryl    Since Andy promoted me to assistant regional manager, I've been tr...  
## 10 Andy     You now, Darryl, this is textbook assistant regional manager stuff...
```

5) str_count()

`str_count()` can be used to count the number of matches in a string

```
theoffice |>
  distinct(character) |>
  slice_head(n = 5) |>
  mutate(
    name = str_to_lower(character), # another way to avoid case sensitivity
    m_s = str_count(name, "m"),
    i_s = str_count(name, "i")
  )
```

```
## # A tibble: 5 × 4
##   character name      m_s   i_s
##   <chr>      <chr>   <int> <int>
## 1 Michael  michael     1     1
## 2 Jim      jim        1     1
## 3 Pam      pam         1     0
## 4 Dwight   dwight      0     1
## 5 Jan      jan         0     0
```

5) str_count() with regex

```
theoffice |>
  distinct(character) |>
  slice_head(n = 5) |>
  mutate(
    name = str_to_lower(character),
    vowels = str_count(name, "[aeiou]"), # count matches of ANY of these characters
    consonants = str_count(name, "[^aeiou]") # count matches of everything EXCEPT these characters
  )
```

```
## # A tibble: 5 × 4
##   character name      vowels consonants
##   <chr>      <chr>    <int>      <int>
## 1 Michael   michael      3          4
## 2 Jim       jim          1          2
## 3 Pam       pam          1          2
## 4 Dwight    dwight       1          5
## 5 Jan       jan          1          2
```

Reminder: `[]` lets you match a set of characters; `^` inverts character class matches

6) str_replace()

As the name suggests, `str_replace()` can be used to modify patterns in strings

```
names
```

```
## [1] "Michael" "Jim" "Pam" "Dwight" "Jan"
## [6] "Michel" "Todd Packer" "Phyllis" "Stanley" "Oscar"
```

```
str_replace(names, "Dw", "Duhw") # jim's office pronunciation guide
```

```
## [1] "Michael" "Jim" "Pam" "Duhwright" "Jan"
## [6] "Michel" "Todd Packer" "Phyllis" "Stanley" "Oscar"
```

6) `str_replace()`

`str_replace()` replaces the first match of a pattern

`str_replace_all()` replaces all matches of a pattern

These functions pair naturally with `mutate()` just like `str_c()`, `str_glue()`, and `str_count()`