

# Time

Week 9

AEM 2850 / 5850 : R for Business Analytics

Cornell Dyson  
Spring 2023

Acknowledgements: Andrew Heiss

# Announcements

Reminders:

- **Group project groups due this Friday (March 24) (canvas link)**
  - Form groups of 3
  - All students in each group must be in the same section (2850 or 5850)
  - If unable to form a group of 3, please state that in your submission
- **Group project due April 14** (full details to come in the next week)

Questions before we get started?

# Plan for today

Prologue

Visualizing time

- Visualizing amounts over time
- Visualizing distributions over time
- Animating plots over time
- Faceting plots over time
- Connecting scatter plots over time

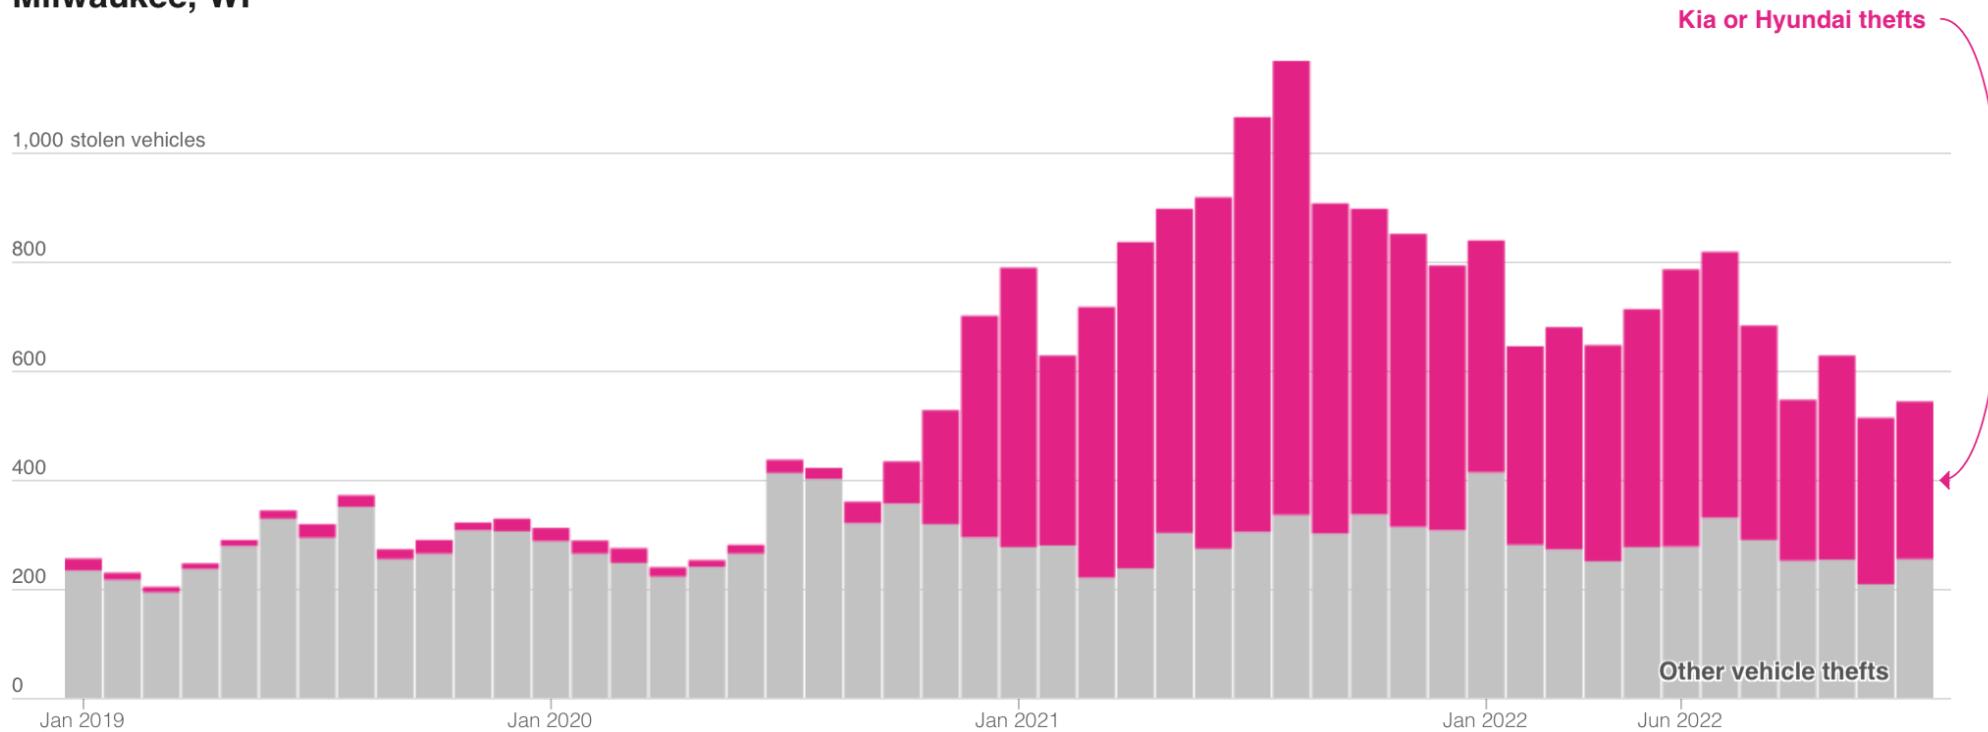
Starting, ending, and decomposing time

Dates and times in R

# Prologue: In the news

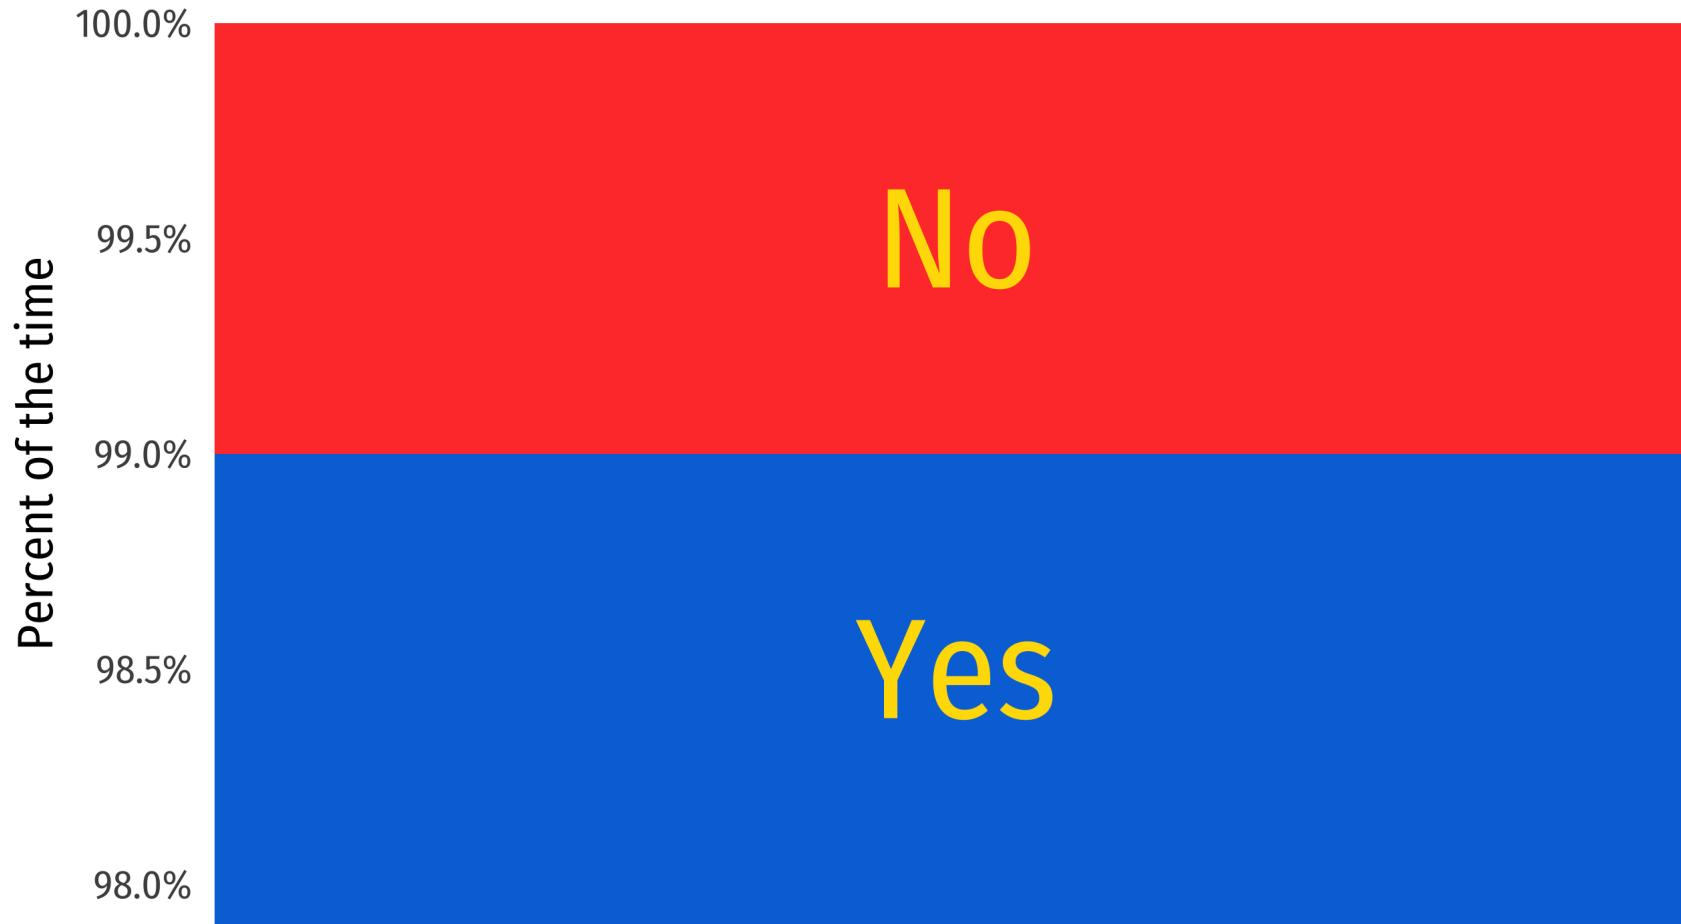
# Car thefts are rising. Is TikTok to blame?

Milwaukee, WI



# Prologue: Axis issues (revisited)

# Is truncating the y-axis misleading?



# Don't be too extreme!

When discussing amounts, we said to never truncate the y-axis

But it is actually more legal to truncate the y-axis than you might think

When do you think it is okay to truncate?

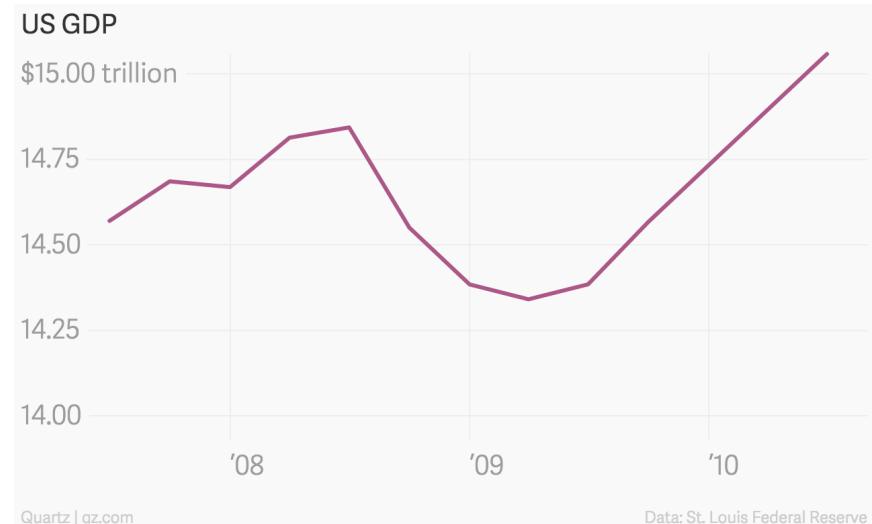
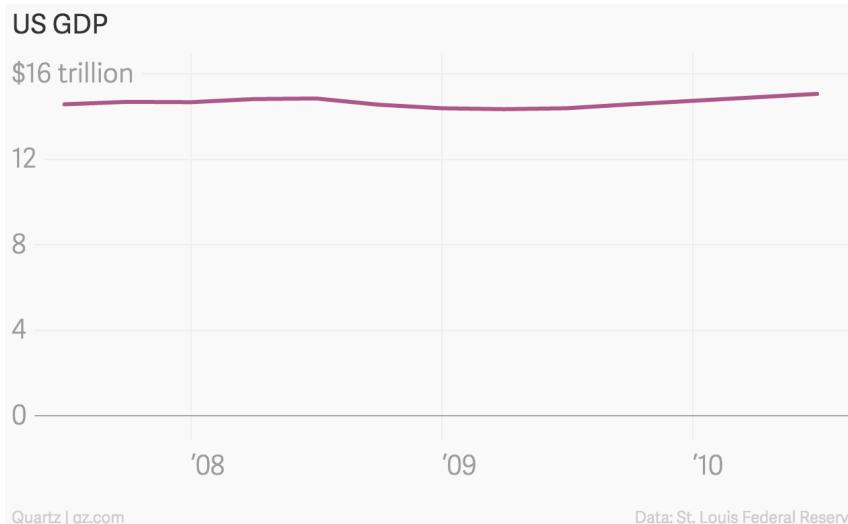
When small movements matter

When the scale itself is distorted

When zero values are impossible

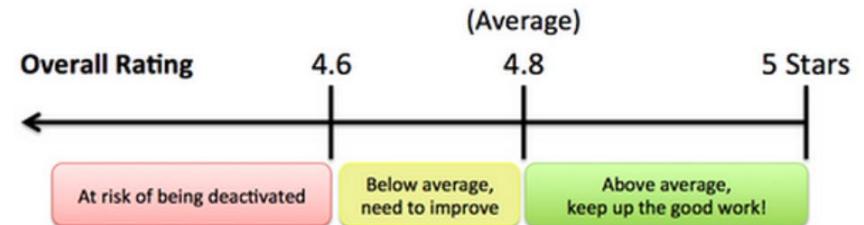
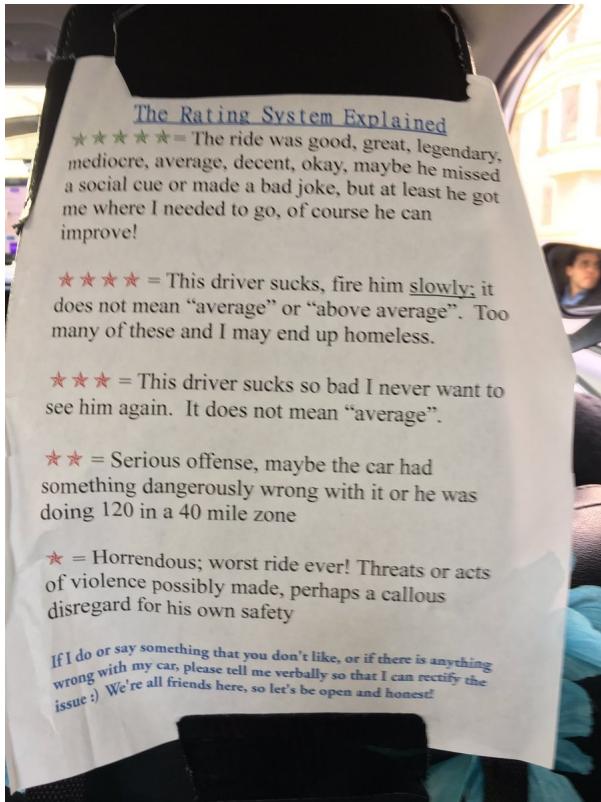
# When is it okay to truncate?

When small movements matter



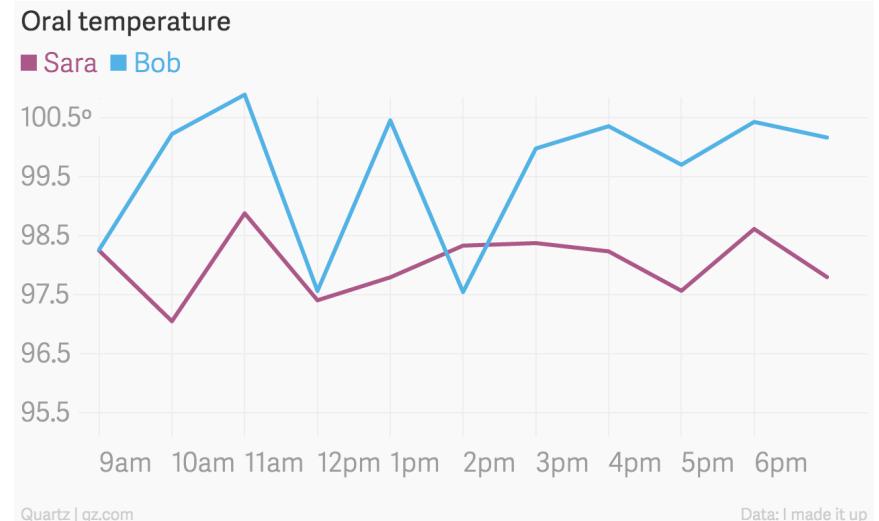
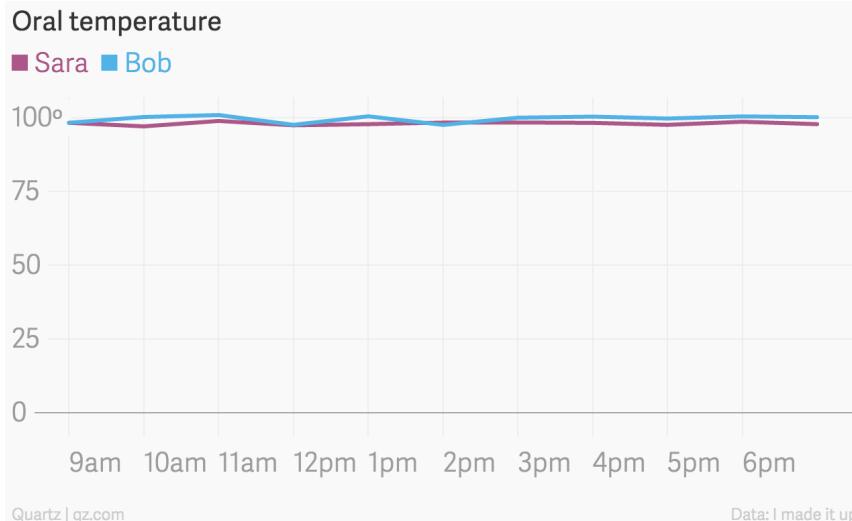
# When is it okay to truncate?

When the scale itself is distorted

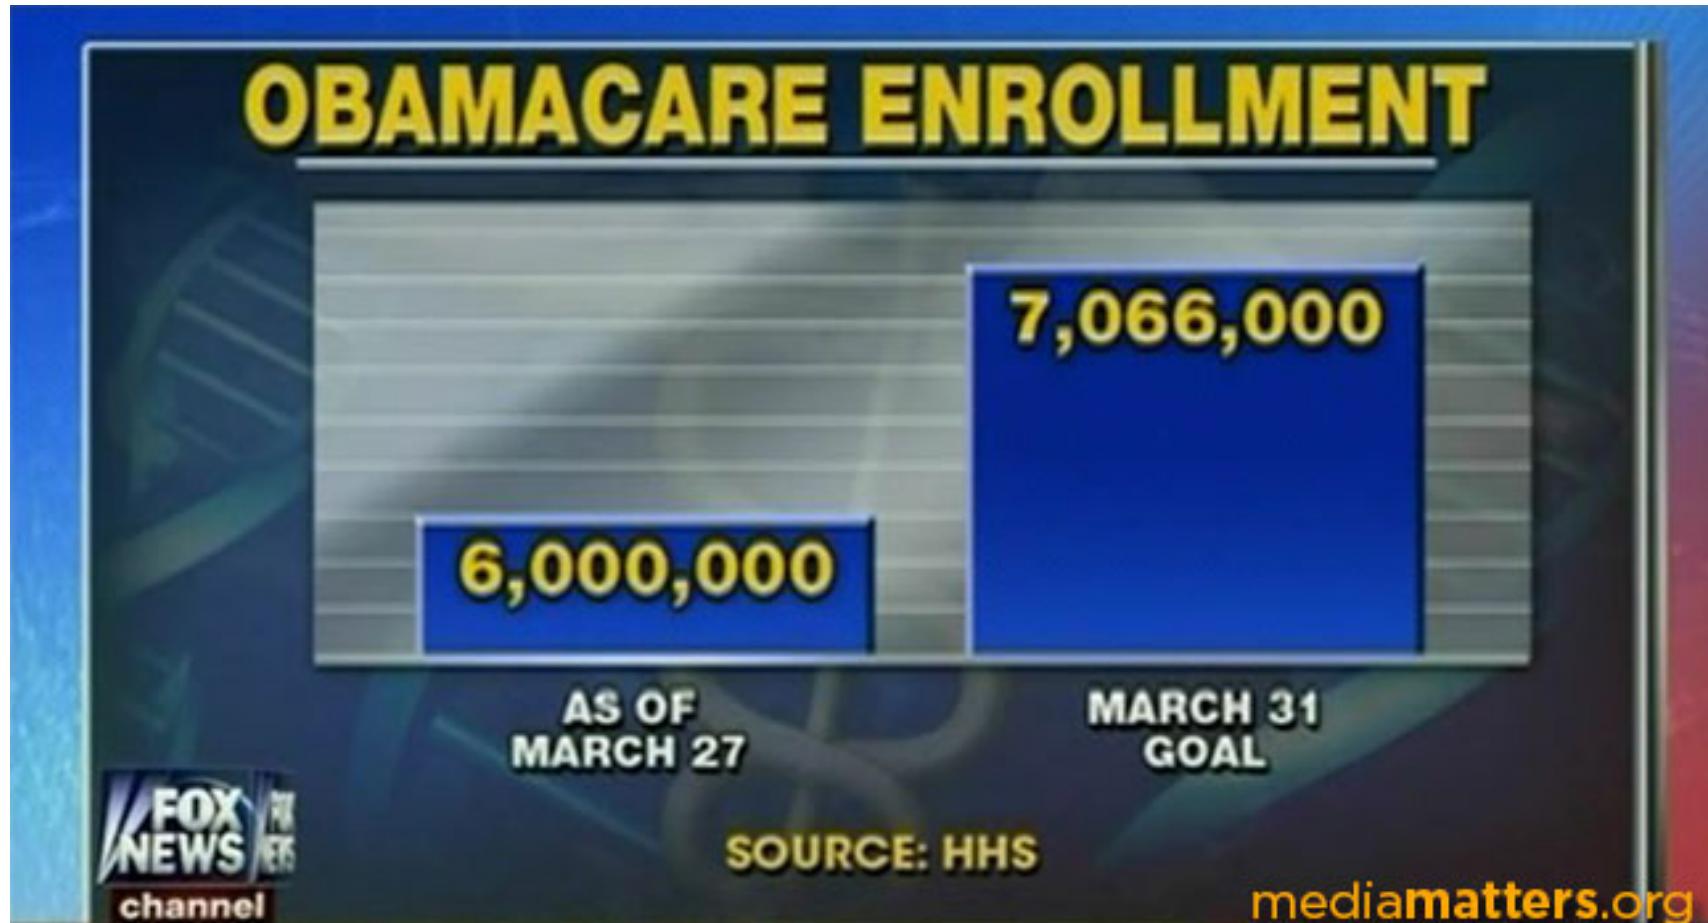


# When is it okay to truncate?

When zero values are impossible



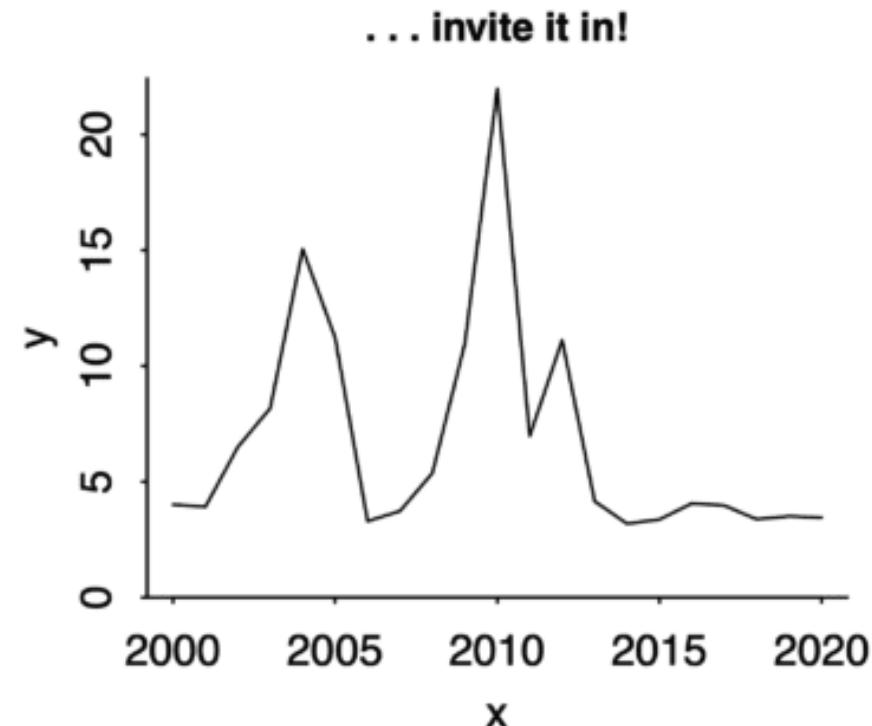
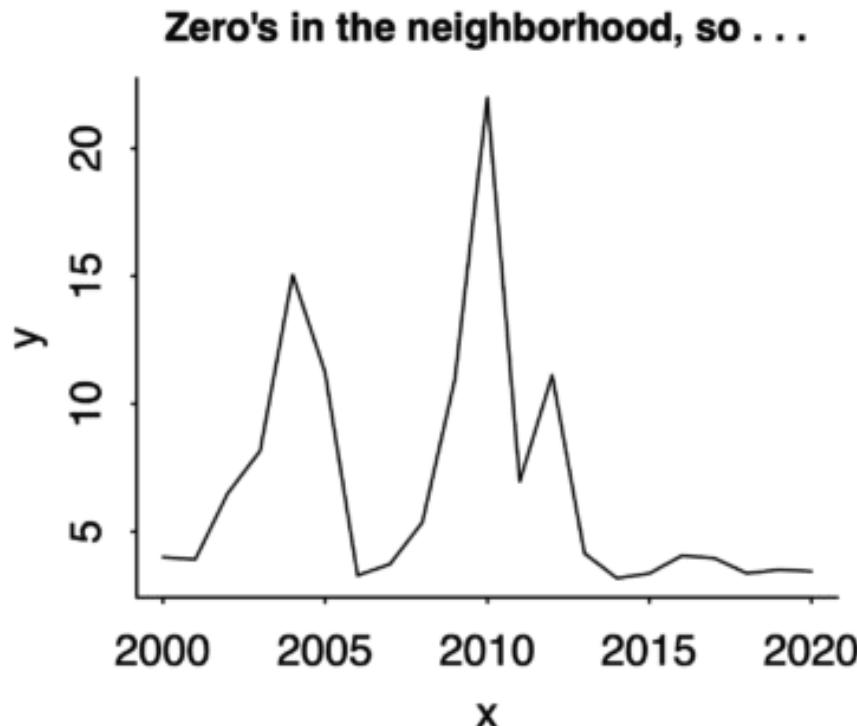
# Never on bar charts



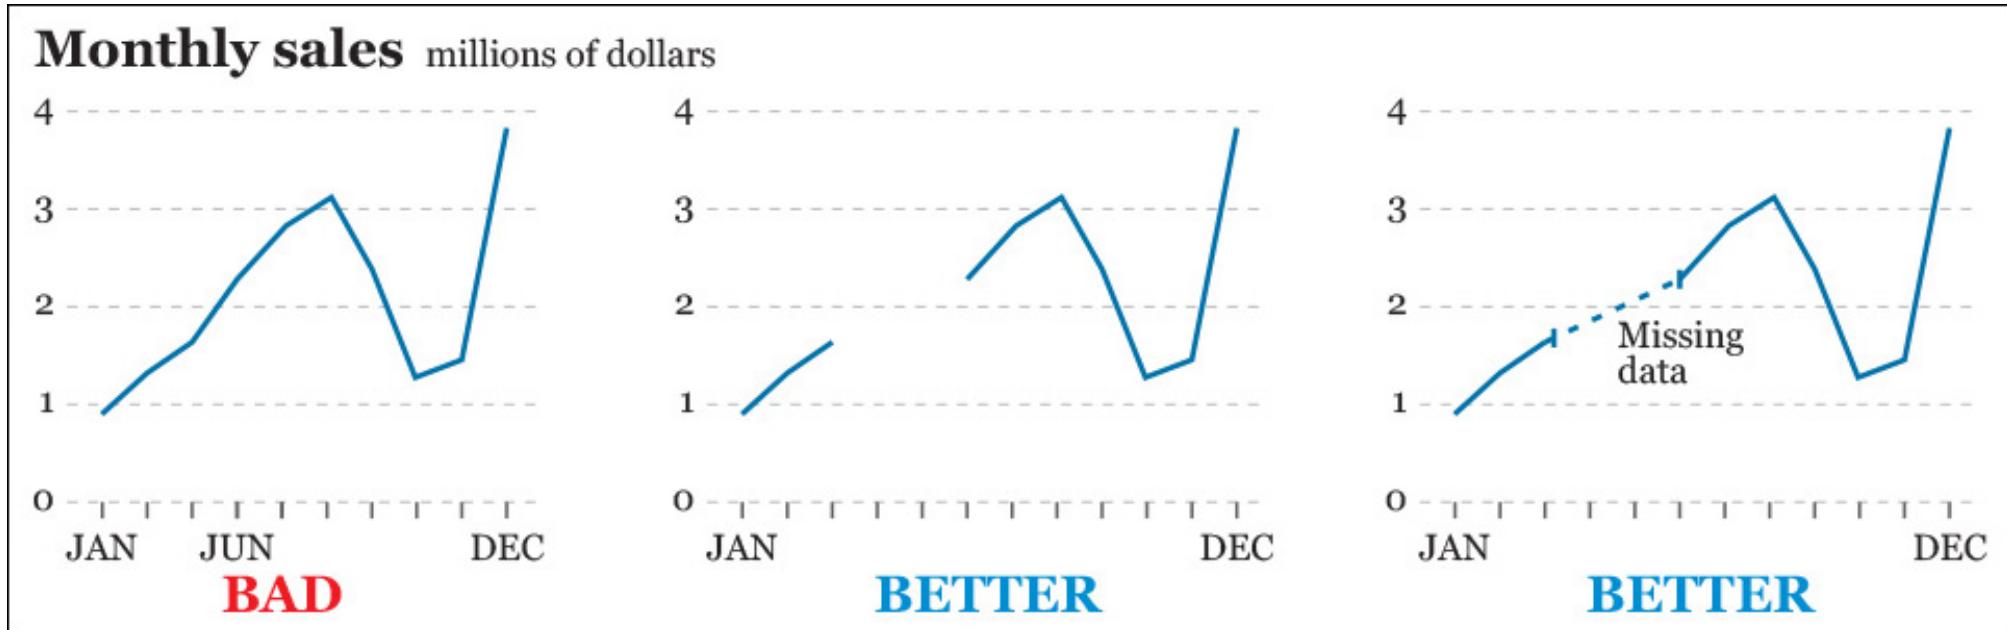
# Zero is okay too!

Just because you don't *have to* start at 0 doesn't mean you should *never* start at 0

Andrew Gelman's heuristic: "**If zero is in the neighborhood, invite it in.**"

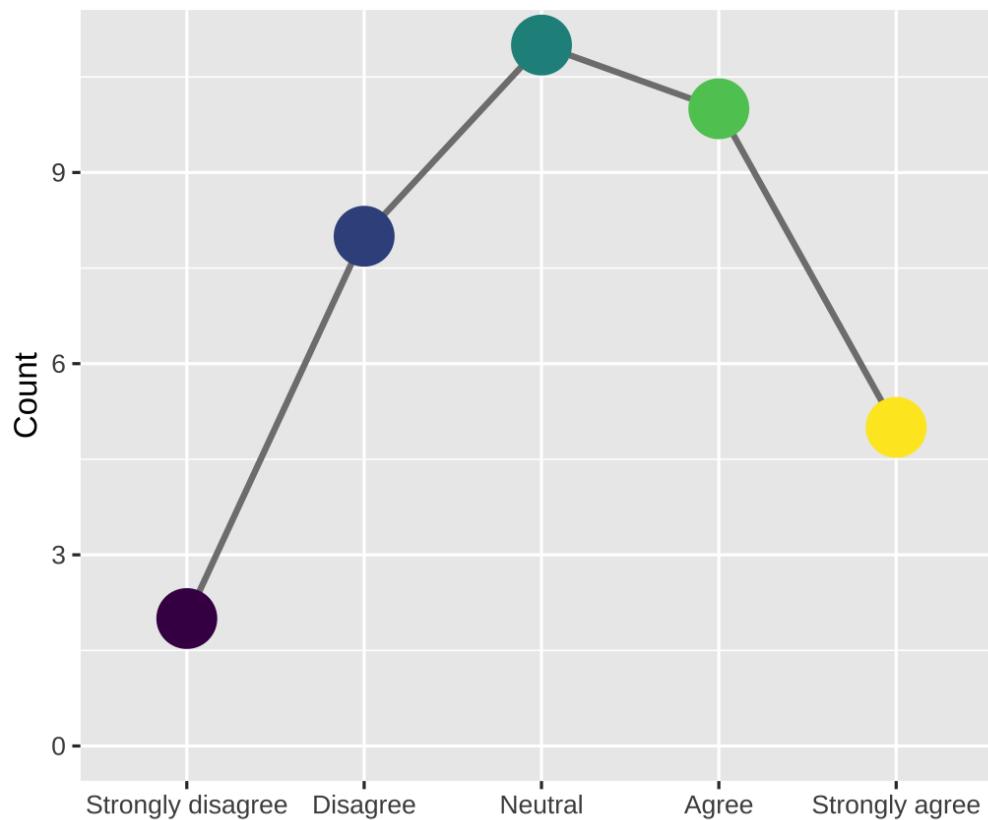


# Also: keep scales consistent, flag missing data

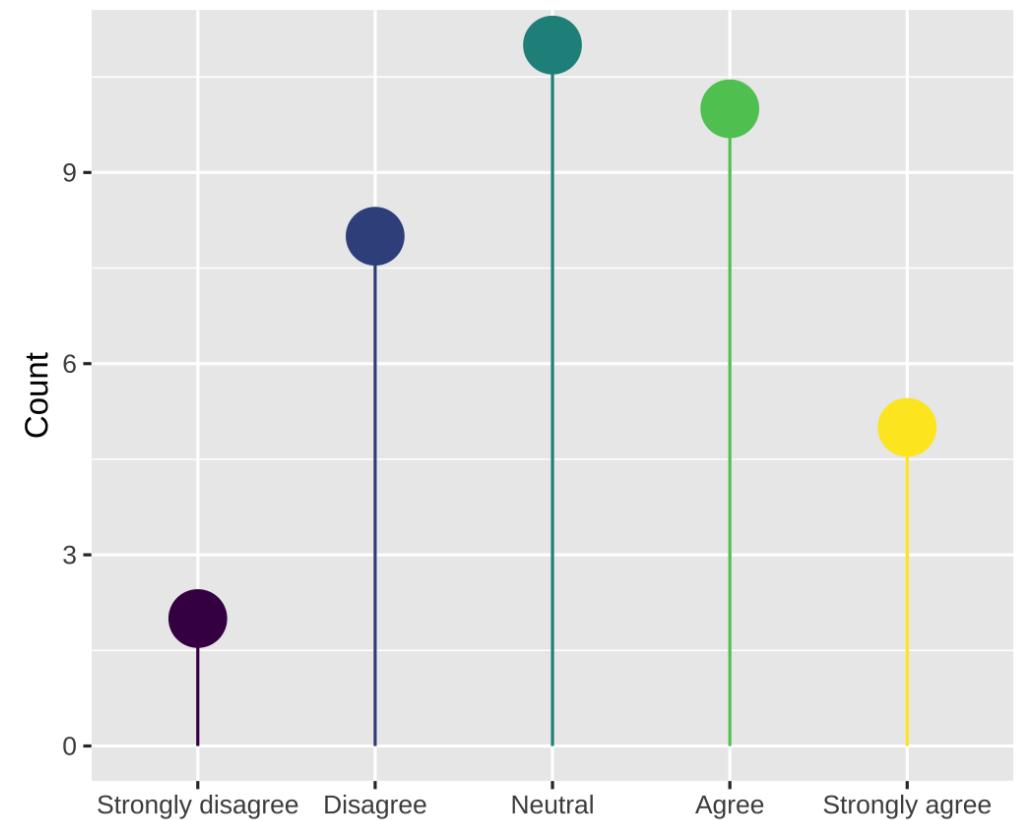


# Also: don't impute across categories

This is BAD



This is BETTER



# Visualizing time

# Visualizing amounts over time

Time is just a variable that can be mapped to an aesthetic

Can be used as `x`, `y`, `color`, `fill`, `facet`, and even animation

Can use all sorts of `geom`s: lines, columns, points, heatmaps, densities, maps, etc.

In general, follow reading conventions to show time progression:

→ & ↓

# Visualizing amounts over time using ggplot

Let's use GameStop share prices to make some plots:

```
gme_prices

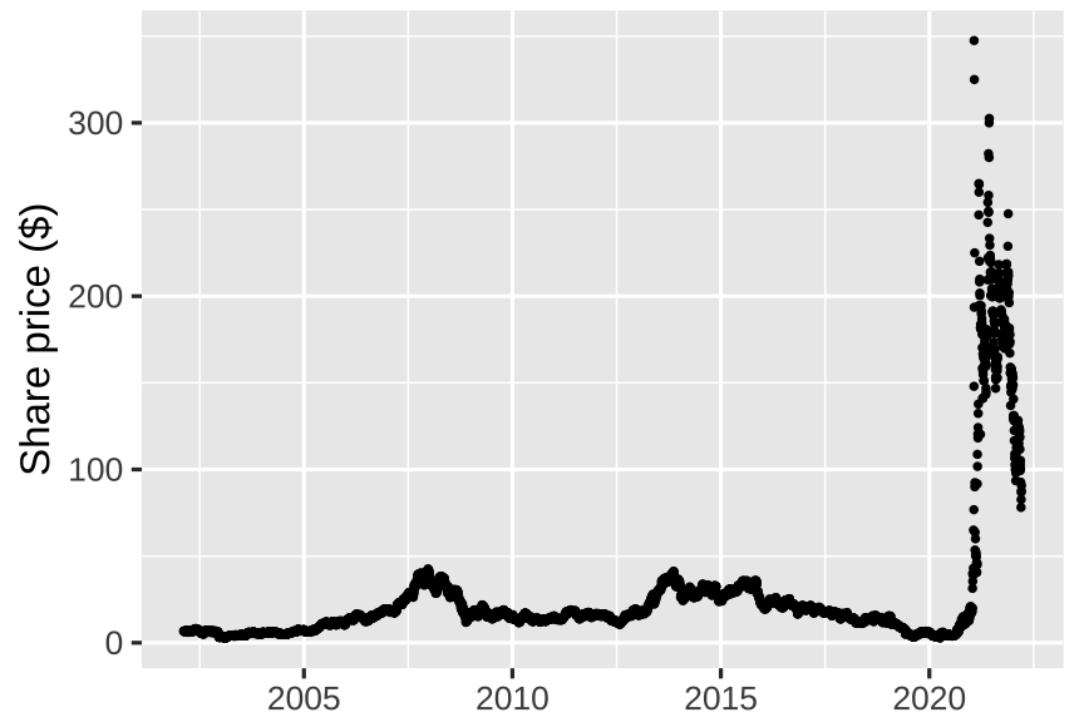
## # A tibble: 5,060 × 8
##   symbol date      open   high    low close volume adjusted
##   <chr>   <date>    <dbl>  <dbl>  <dbl>  <dbl>    <dbl>    <dbl>
## 1 GME    2002-02-13 9.62  10.1   9.52  10.0   19054000  6.77
## 2 GME    2002-02-14 10.2   10.2   9.93  10     2755400   6.73
## 3 GME    2002-02-15 10     10.0   9.85  9.95   2097400   6.70
## 4 GME    2002-02-19 9.9    9.9    9.38  9.55   1852600   6.43
## 5 GME    2002-02-20 9.6    9.88   9.52  9.88   1723200   6.65
## 6 GME    2002-02-21 9.84   9.93   9.75  9.85   1744200   6.63
## 7 GME    2002-02-22 9.93   9.93   9.6   9.68   881400    6.51
## 8 GME    2002-02-25 9.65   9.82   9.54  9.75   863400    6.56
## 9 GME    2002-02-26 9.7    9.85   9.54  9.75   690400    6.56
## 10 GME   2002-02-27 9.68   9.68   9.5   9.57   1022800   6.45
## # ... with 5,050 more rows
```

# Time on x-axis + geom\_point()

```
gme_prices |>  
  ggplot(aes(x = date, y = adjusted)) +  
  geom_point(size = 0.5) +  
  labs(x = NULL, y = "Share price ($)",  
       title = "GME to the moon")
```

How would you add a line to this?

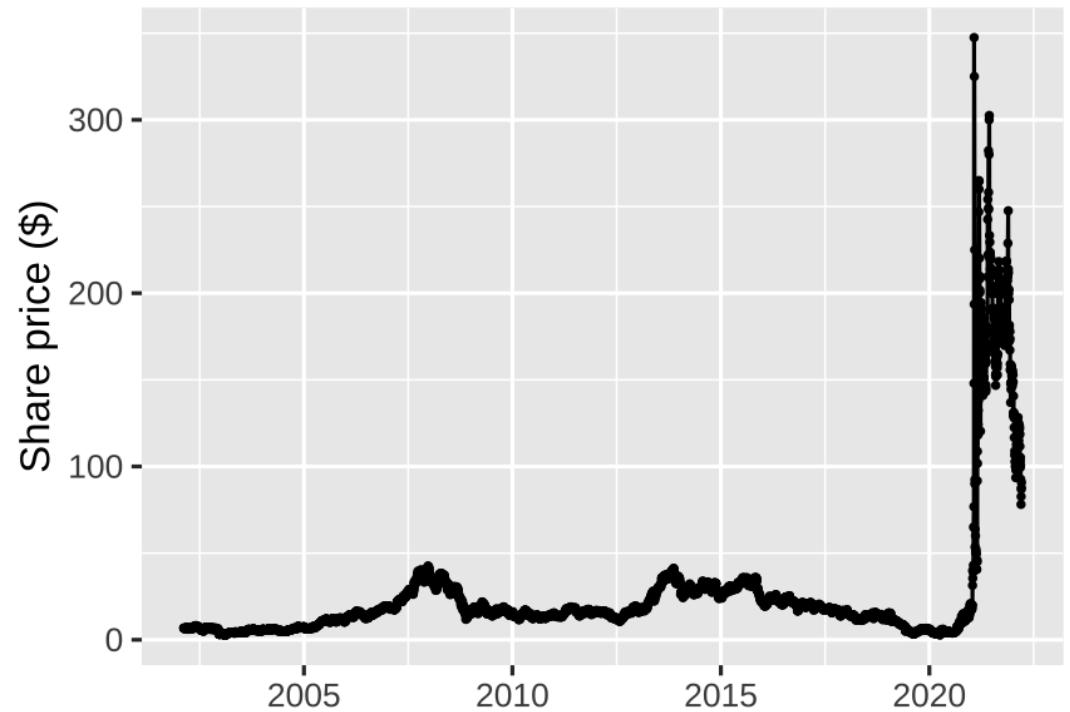
GME to the moon



# Time on x-axis + geom\_point()

```
gme_prices |>  
  ggplot(aes(x = date, y = adjusted)) +  
  geom_point(size = 0.5) +  
  geom_line() +  
  labs(x = NULL, y = "Share price ($)",  
       title = "GME to the moon")
```

GME to the moon



# Time on x-axis: points vs lines

Points emphasize observations

Lines emphasize trends

Using lines for time series is often fine since data are evenly spaced and usually complete

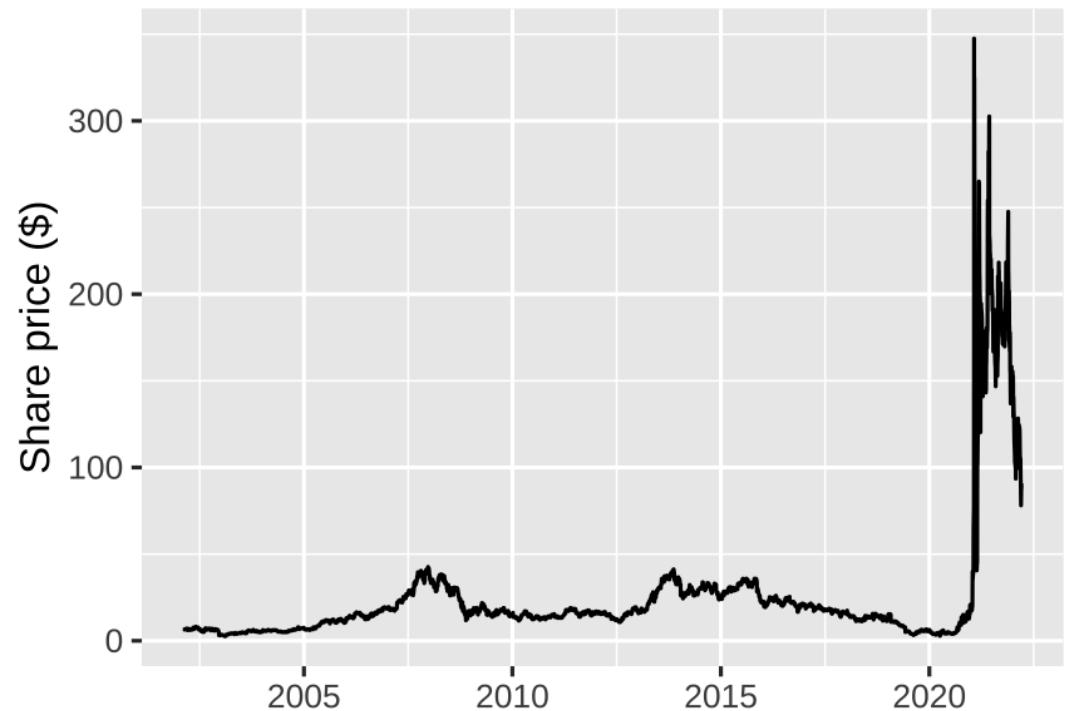
But lines are effectively made-up data, so be careful how you use them!

# Time on x-axis + geom\_line/col()

The GameStop share price plot is clearer without points:

```
gme_prices |>  
  ggplot(aes(x = date, y = adjusted)) +  
  geom_line() +  
  labs(x = NULL, y = "Share price ($)",  
       title = "GME to the moon")
```

GME to the moon

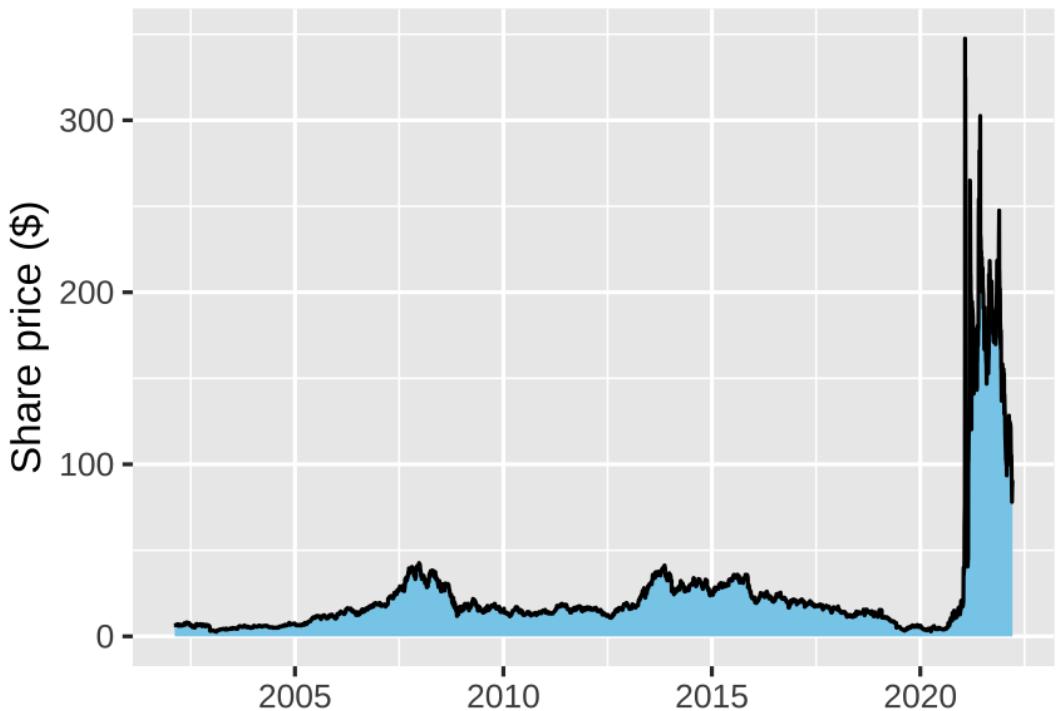


# Time on x-axis + geom\_line/col()

Can also use a fill for the area under a line, **as long as the y-axis starts at zero**

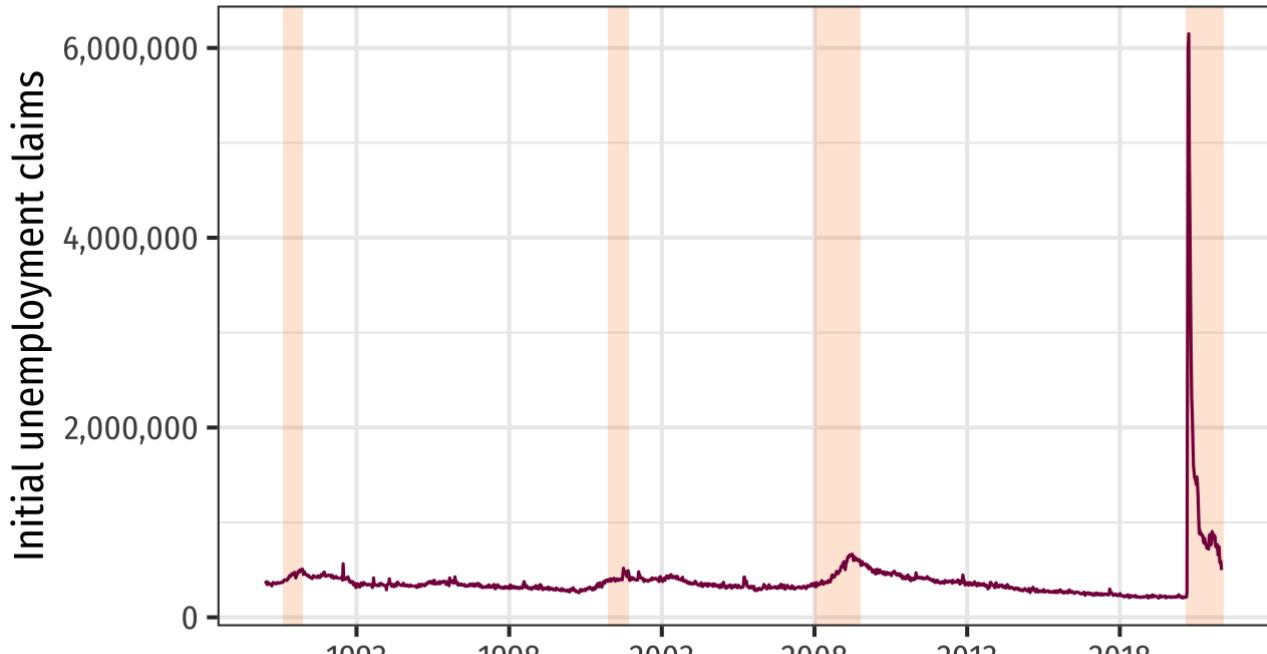
```
gme_prices |>  
  ggplot(aes(x = date, y = adjusted)) +  
  geom_area(fill = "skyblue", color = "black") +  
  labs(x = NULL, y = "Share price ($)",  
       title = "GME to the moon")
```

GME to the moon



# Line plots don't have to be boring

HOLY CRAP



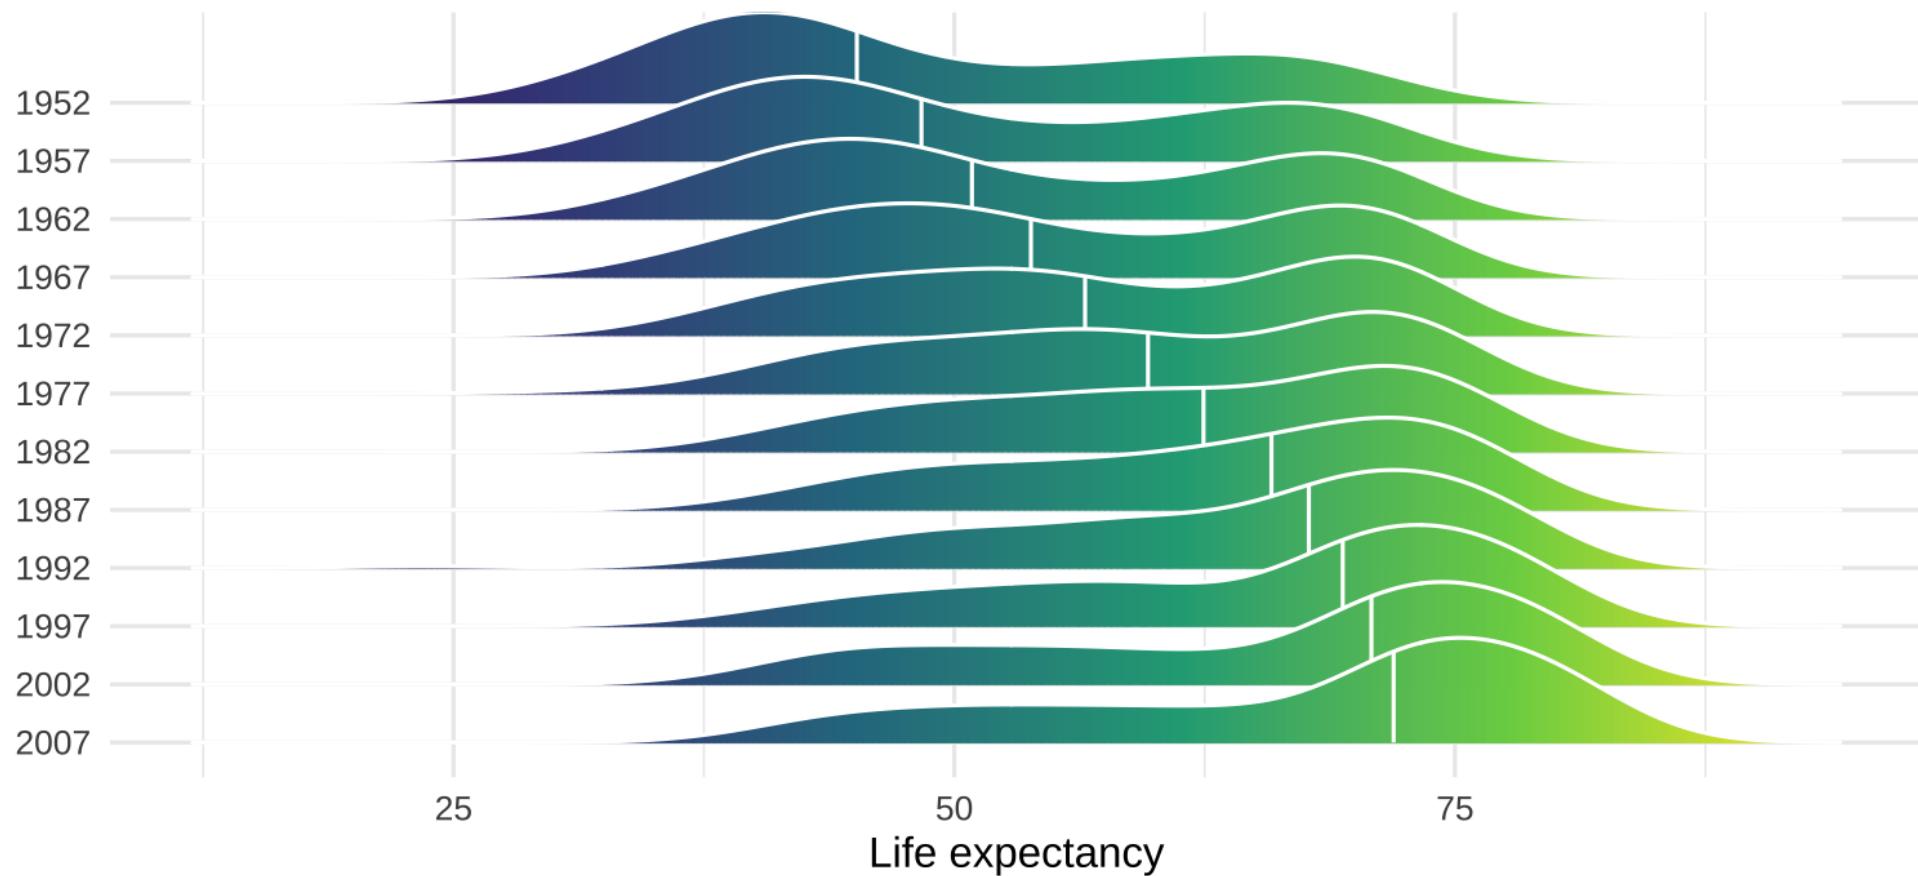
Source: Initial weekly unemployment claims (ICSA); FRED  
Recessions highlighted in orange



Source: U.S. Bureau of Labor Statistics

# Visualizing distributions over time

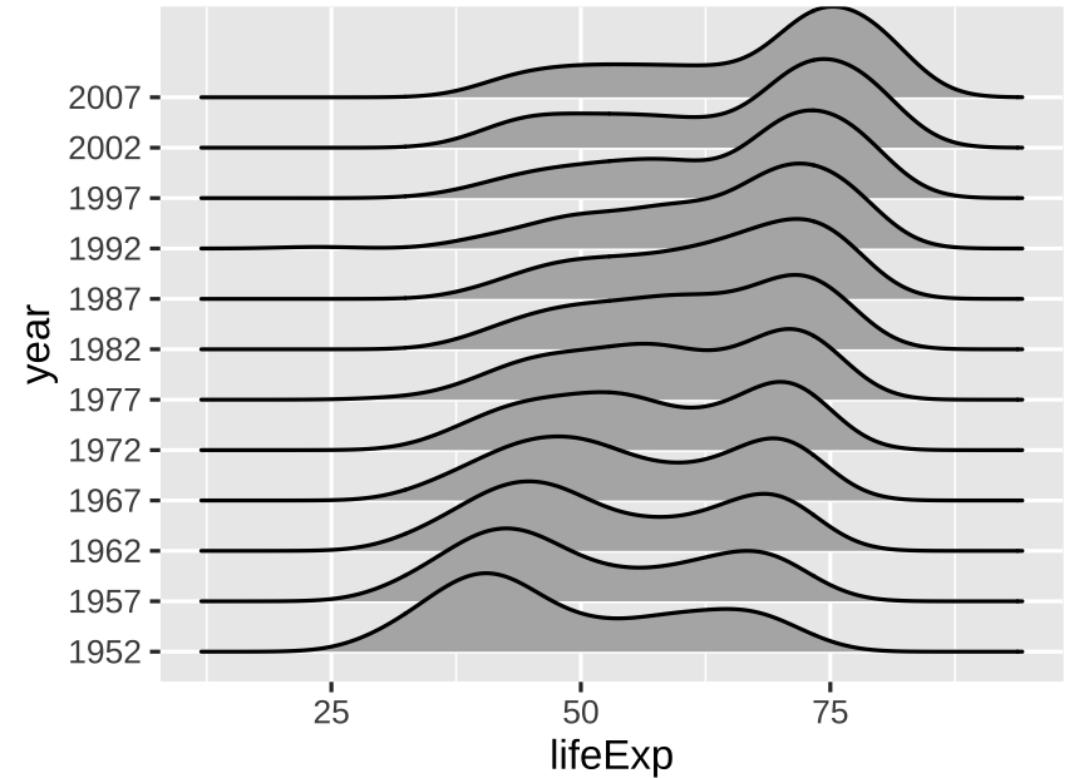
We can also visualize how distributions evolve. How could you make this plot?



# Let's start simple

```
library(ggridges) # for geom_density_ridges()  
  
gapminder |>  
  mutate(year = as_factor(year)) |>  
  ggplot(  
    aes(x = lifeExp,  
        y = year) # map time to y  
  ) +  
  geom_density_ridges() # add geom
```

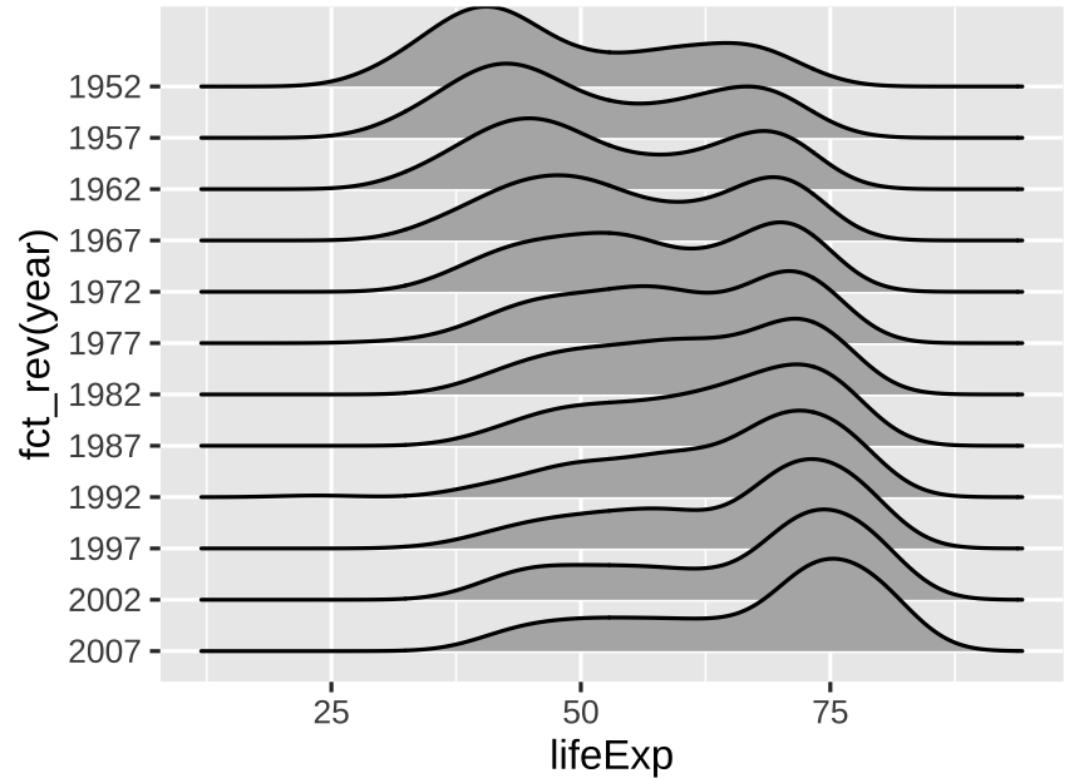
How could we modify this to follow ↓ convention for time?



# Follow ↓ convention for time

```
gapminder |>  
  mutate(year = as_factor(year)) |>  
  ggplot(  
    aes(x = lifeExp,  
        y = fct_rev(year)) # reverse year  
  ) +  
  geom_density_ridges()
```

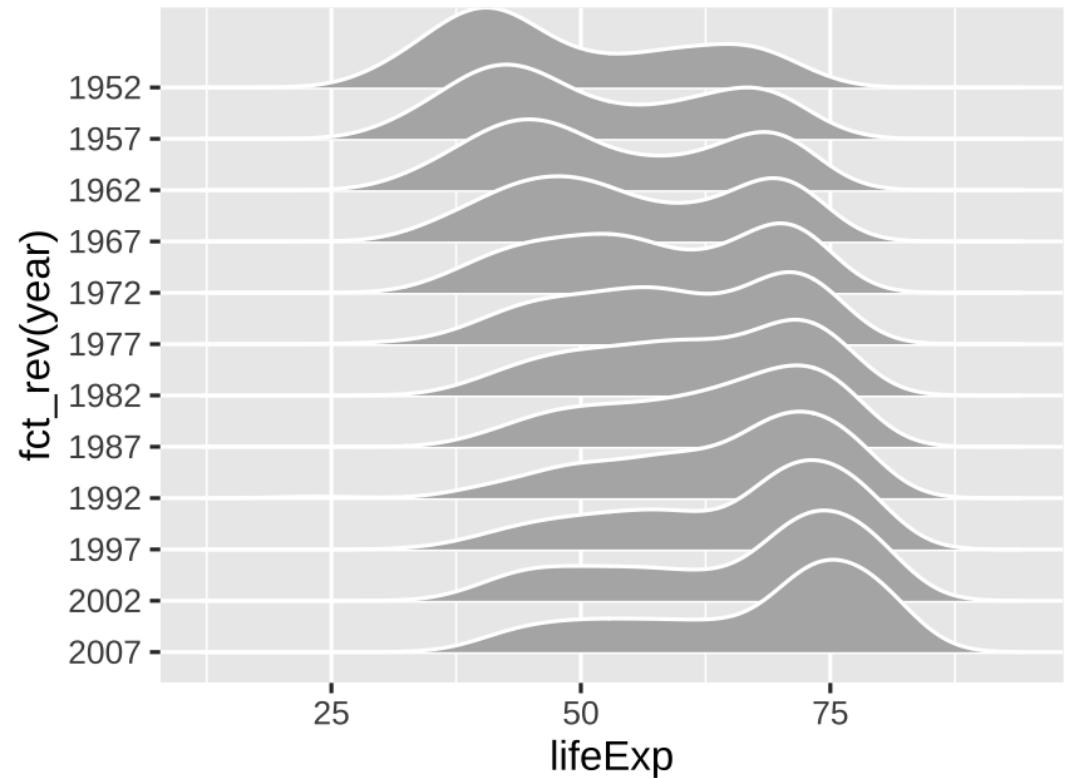
How could we make the lines white?



# Add color to help visually separate densities

```
gapminder |>  
  mutate(year = as_factor(year)) |>  
  ggplot(  
    aes(x = lifeExp,  
        y = fct_rev(year))  
  ) +  
  geom_density_ridges(  
    color = "white" # separate densities  
  )
```

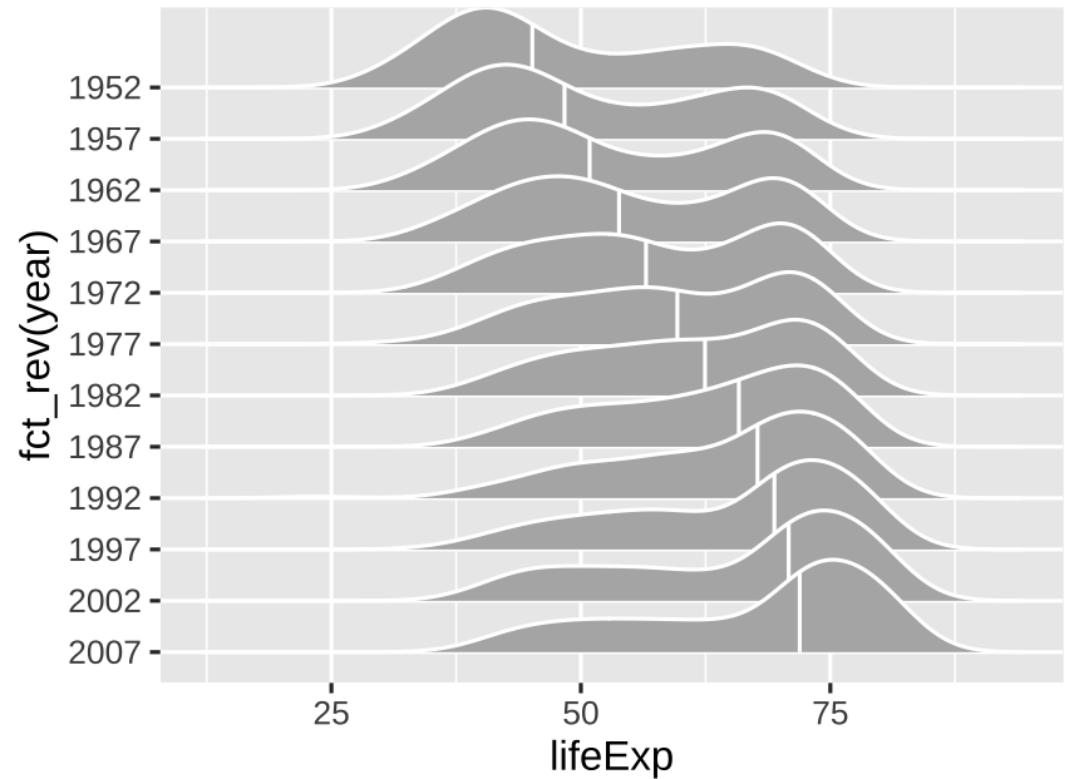
How could we add a line at the median of each density?



# Add a line at the median of each density

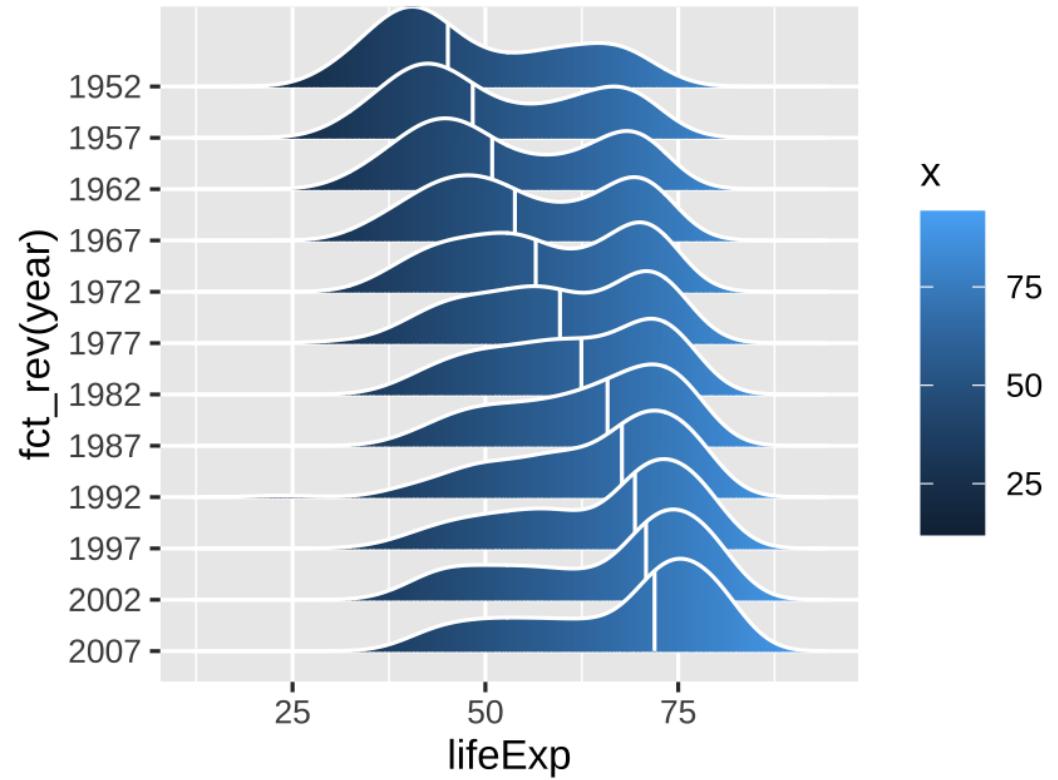
```
gapminder |>  
  mutate(year = as_factor(year)) |>  
  ggplot(  
    aes(x = lifeExp,  
        y = fct_rev(year))  
  ) +  
  geom_density_ridges(  
    color = "white",  
    quantile_lines = TRUE, # add lines  
    quantiles = 2           # median  
  )
```

How could we add a color gradient to the densities?



# Fill the densities along the x axis

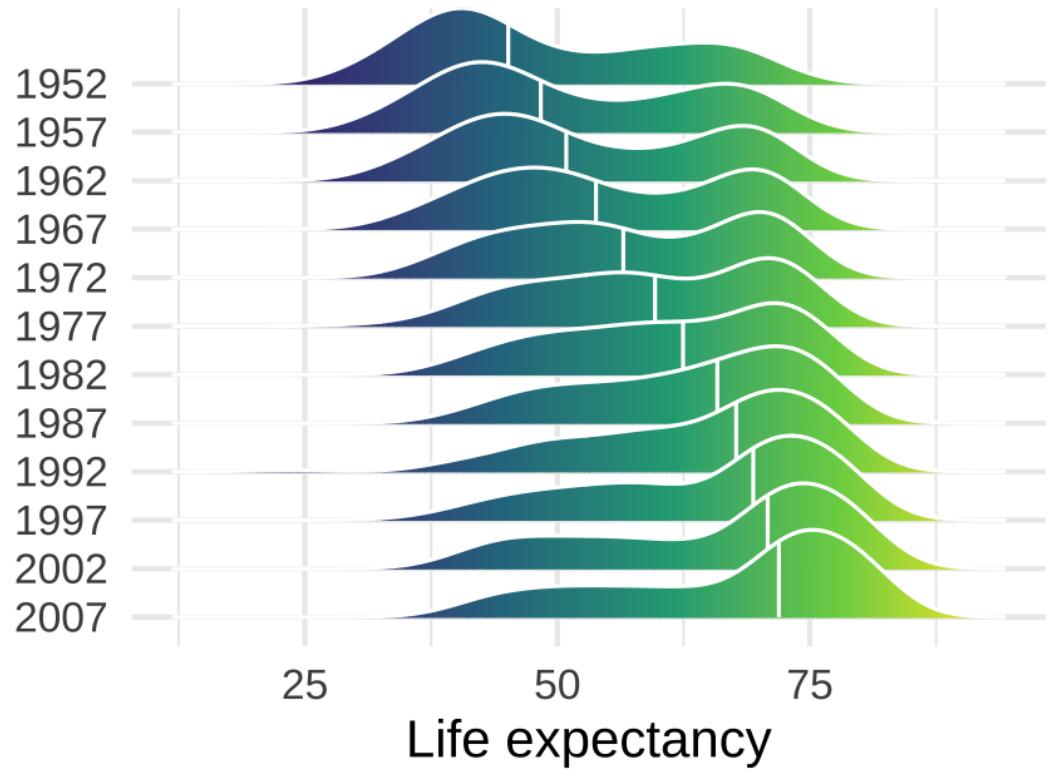
```
gapminder |>  
  mutate(year = as_factor(year)) |>  
  ggplot(  
    aes(x = lifeExp,  
        y = fct_rev(year),  
        fill = after_stat(x) # fill by x  
    ) +  
  geom_density_ridges_gradient(  
    # this is a special geom to fill along x  
    color = "white",  
    quantile_lines = TRUE,  
    quantiles = 2  
  )
```



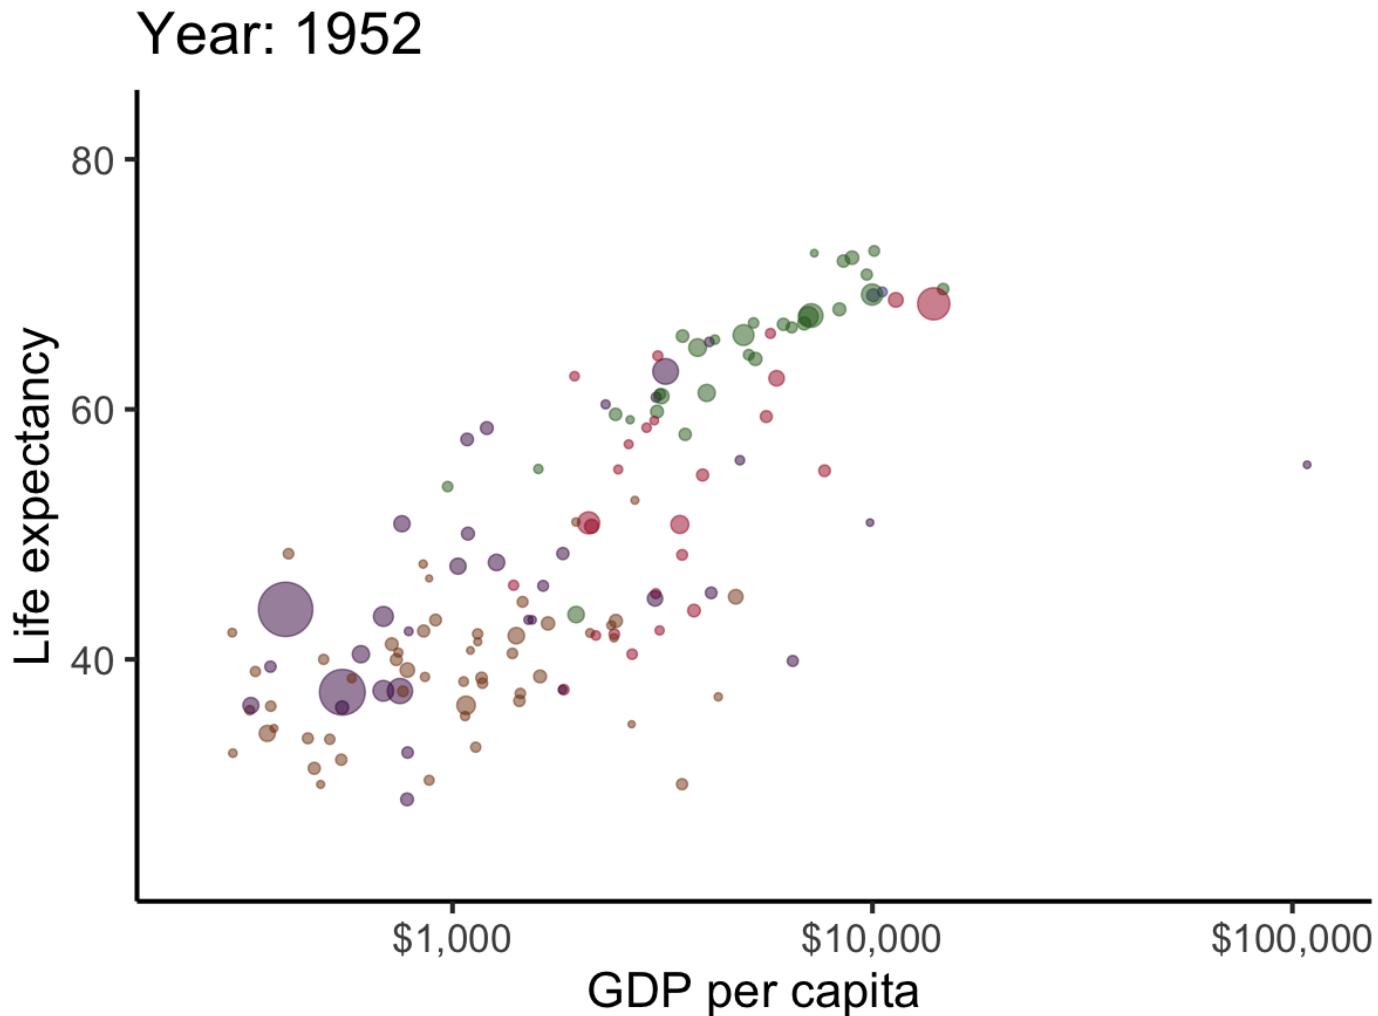
Use `after_stat(x)` to shade the *densities*, not the data itself

# Finish by cleaning the plot up

```
gapminder |>  
  mutate(year = as_factor(year)) |>  
  ggplot(  
    aes(x = lifeExp,  
        y = fct_rev(year),  
        fill = after_stat(x))  
  ) +  
  geom_density_ridges_gradient(  
    color = "white",  
    quantile_lines = TRUE,  
    quantiles = 2  
  ) +  
  guides(fill = "none") +      # omit legend  
  scale_fill_viridis_c() +     # nicer shading  
  labs(x = "Life expectancy",# modify label  
       y = NULL,             # omit labels  
       fill = NULL) +        # omit labels  
  theme_minimal(              # cleaner theme  
    base_size = 14          # larger font  
  )
```



# Animating plots over time

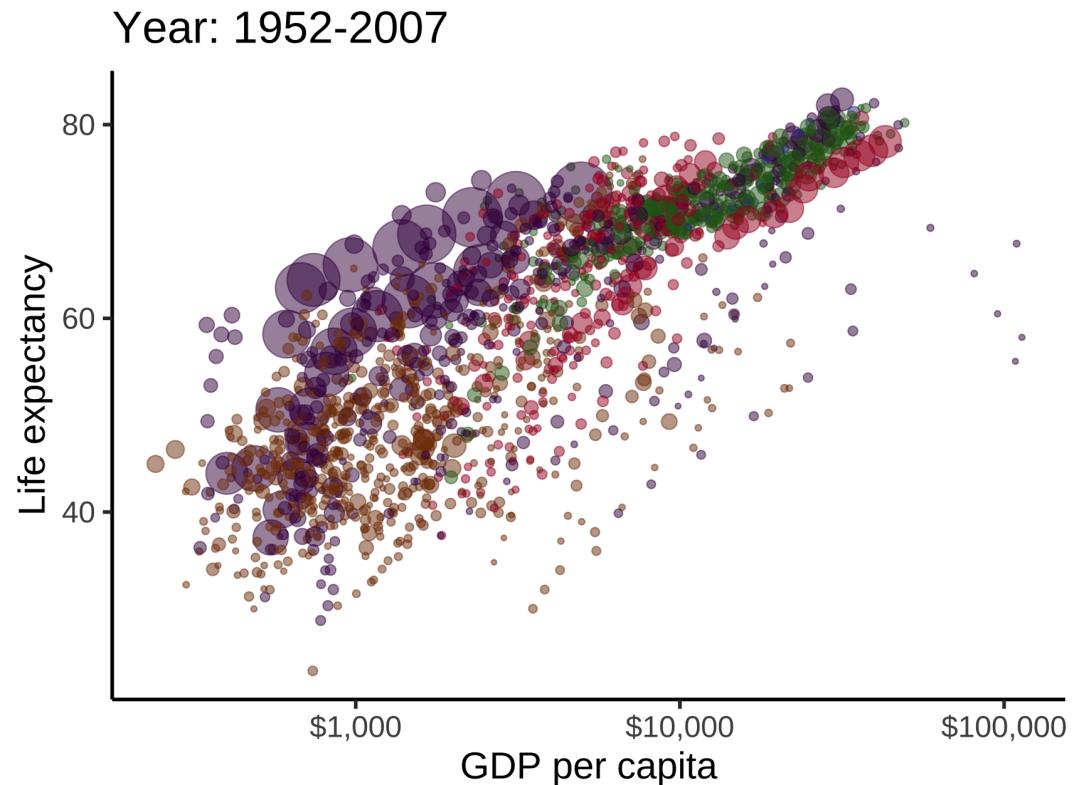


# How can we make this animation in R?

First, how would we make a static plot of the data?

```
library(gapminder) # gapminder data

gapminder |>
  ggplot(aes(gdpPercap, lifeExp,
             size = pop, color = continent)) +
  geom_point(alpha = 0.5, show.legend = FALSE)
  scale_color_manual(
    values = gapminder::continent_colors
  ) +
  scale_size_continuous(range = c(1, 15)) +
  scale_x_log10(labels = label_dollar()) +
  theme_classic(base_size = 20) +
  labs(x = "GDP per capita",
       y = "Life expectancy",
       title = "Year: 1952-2007")
```

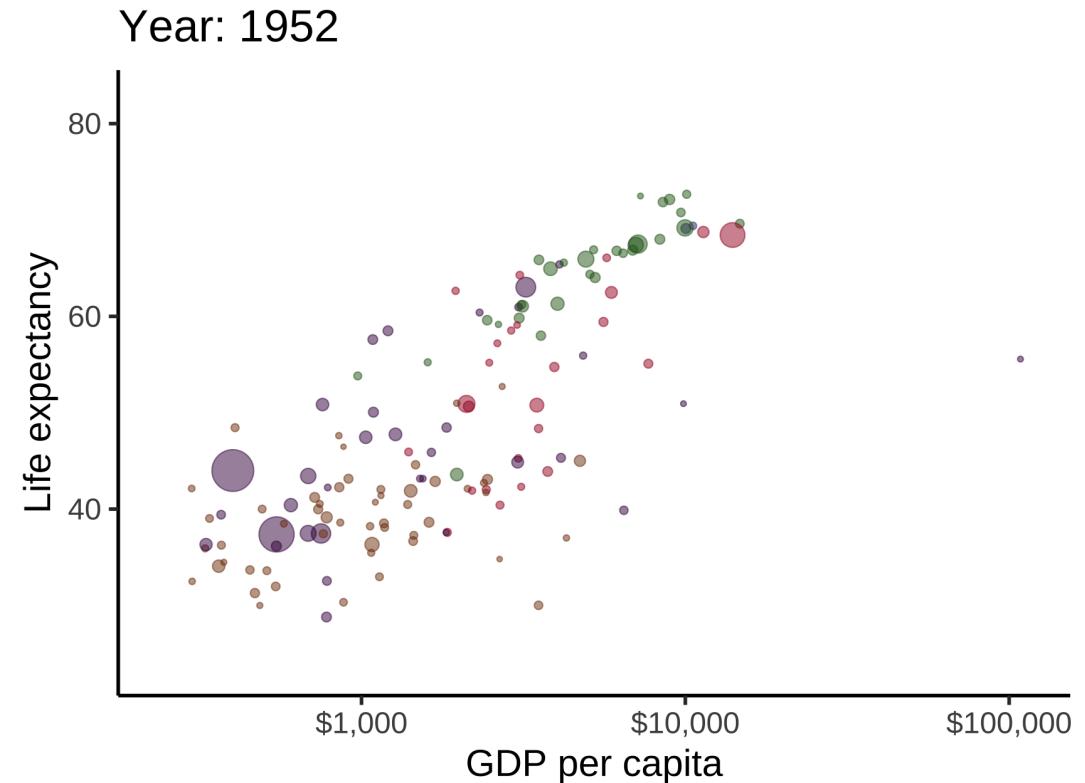


# Converting a static plot to an animation

Second, let's use `ganimate` to visualize changes over time

```
library(gganimate) # plot animation package

gapminder |>
  ggplot(aes(gdpPercap, lifeExp,
             size = pop, color = continent)) +
  geom_point(alpha = 0.5, show.legend = FALSE)
  scale_color_manual(
    values = gapminder::continent_colors
  ) +
  scale_size_continuous(range = c(1, 15)) +
  scale_x_log10(labels = label_dollar()) +
  theme_classic(base_size = 20) +
  labs(x = "GDP per capita",
       y = "Life expectancy",
       title = "Year: {frame_time}") +
  transition_time(year) +
  ease_aes('linear') # default progression
```



# Saving gifs

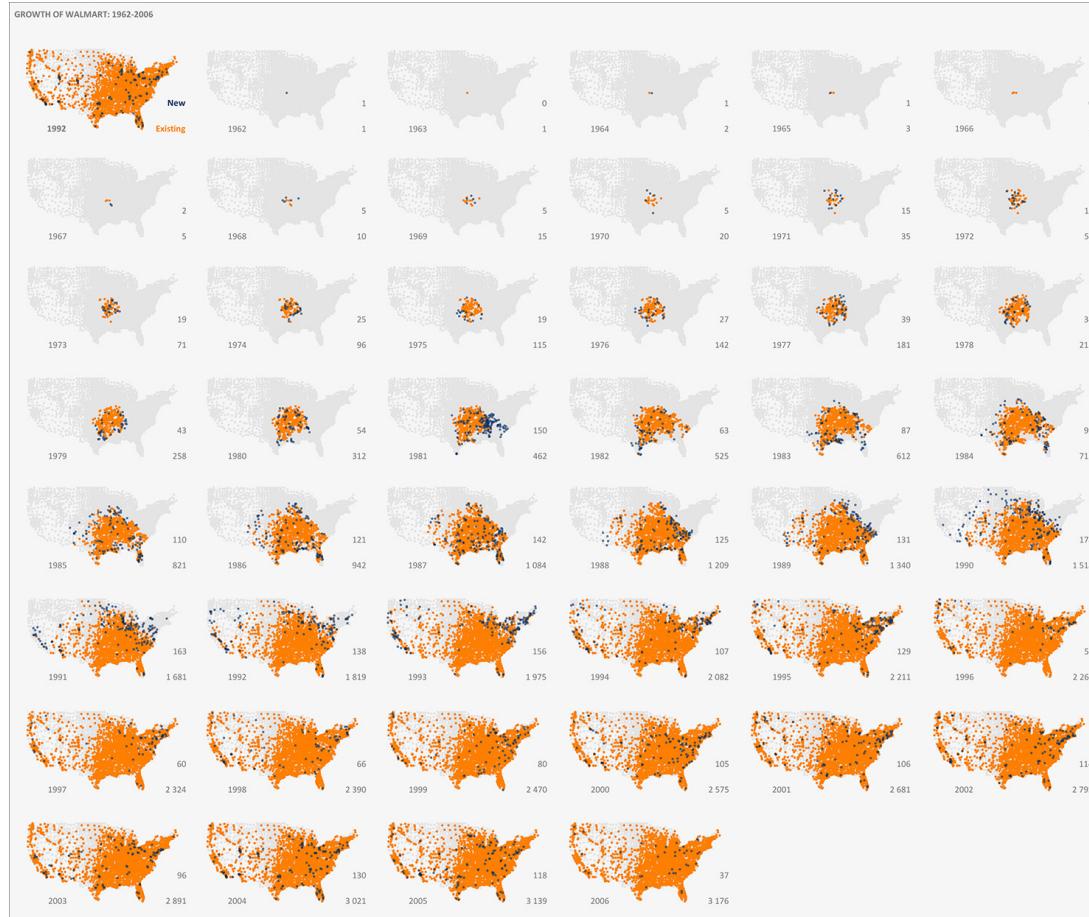
Third, use `anim_save()` to write a `.gif` that you can text to your mom

```
library(gganimate)
library(gifski)

my_gapminder_animation <- gapminder |>
  ggplot(aes(gdpPercap, lifeExp, size = pop, color = continent)) +
  geom_point(alpha = 0.5, show.legend = FALSE) +
  scale_color_manual(values = gapminder::continent_colors) +
  scale_size_continuous(range = c(1, 15)) +
  scale_x_log10(labels = label_dollar()) +
  theme_classic(base_size = 20) +
  labs(x = "GDP per capita",
       y = "Life expectancy",
       title = "Year: {frame_time}") +
  transition_time(year) +
  ease_aes('linear')

animate(my_gapminder_animation, renderer = gifski_renderer())
anim_save("content/slides/img/09/my-gapminder-animation.gif")
```

# Faceting plots over time



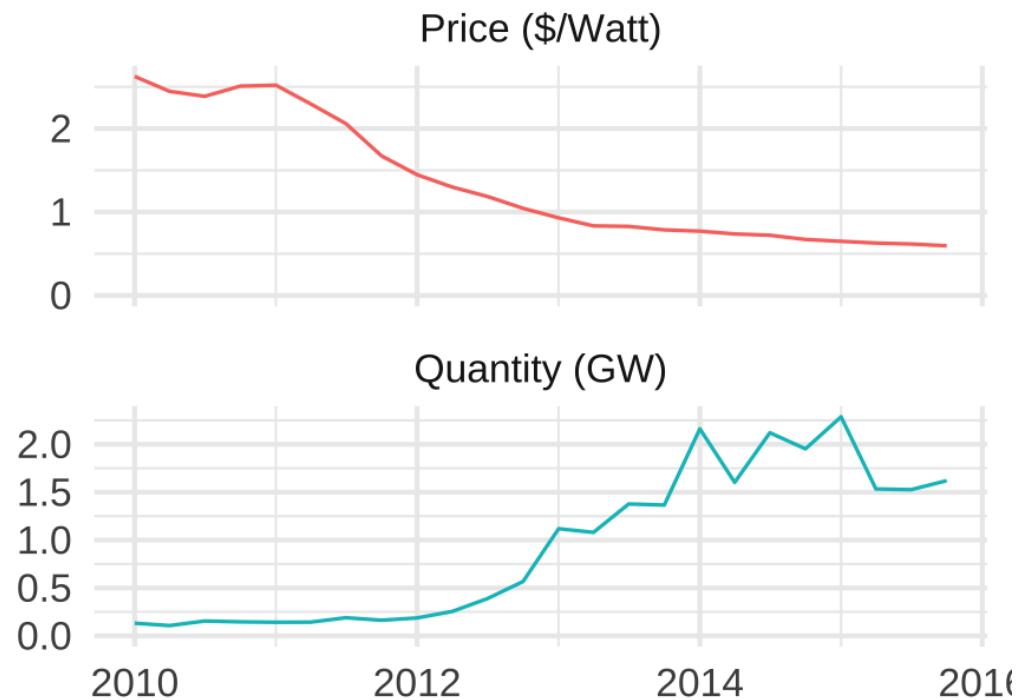
Map of the spread of Walmart by Jorge Camões

# Connecting scatter plots over time

Sometimes connected scatter plots of time series data make sense

What is a connected scatter plot?

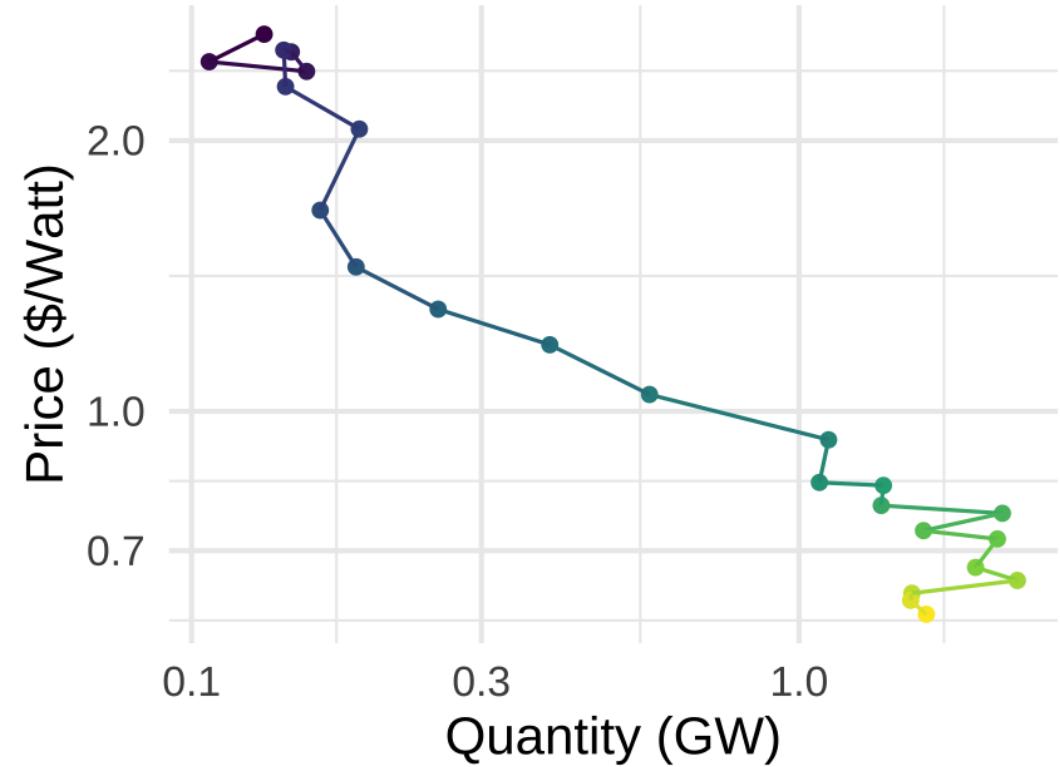
How would you make one using these data?



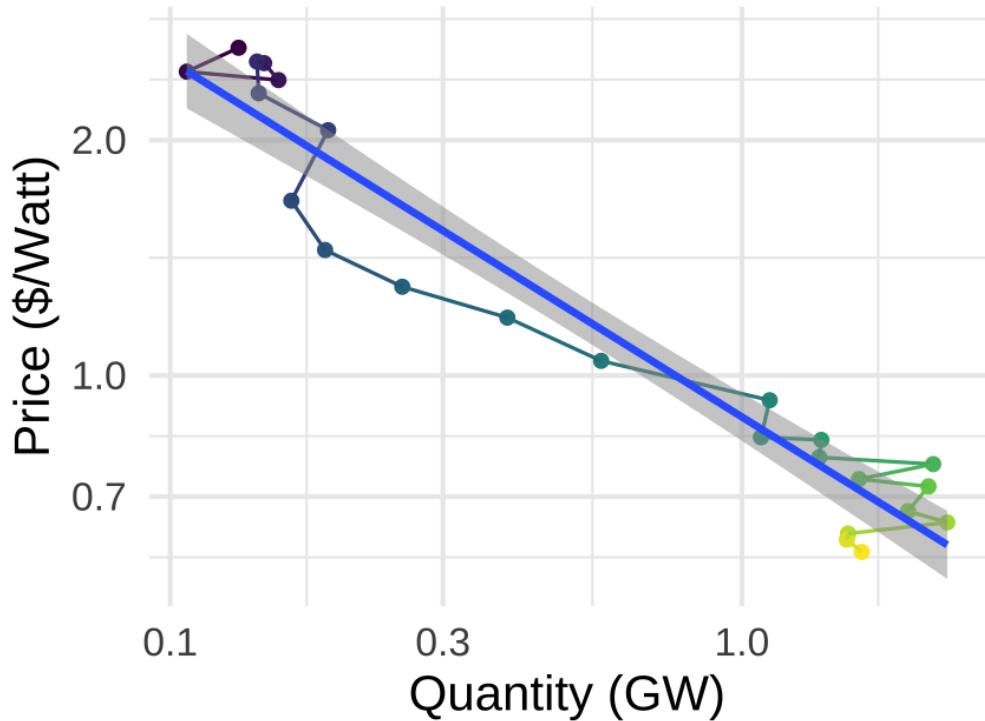
# Connected scatter plots

```
solar_data |>  
  ggplot(aes(x = `Quantity (GW)`, # not time  
             y = `Price ($/Watt)`, # not time  
             color = date)) +      # time!  
  geom_point() +  
  geom_path() + # connect by time, not x  
  scale_x_log10() +  
  scale_y_log10() +  
  scale_color_viridis_c() +  
  guides(color = "none") +  
  theme_minimal(base_size = 14)
```

Note the log axes



# Is this a good use case?



Looks a lot like a demand curve!

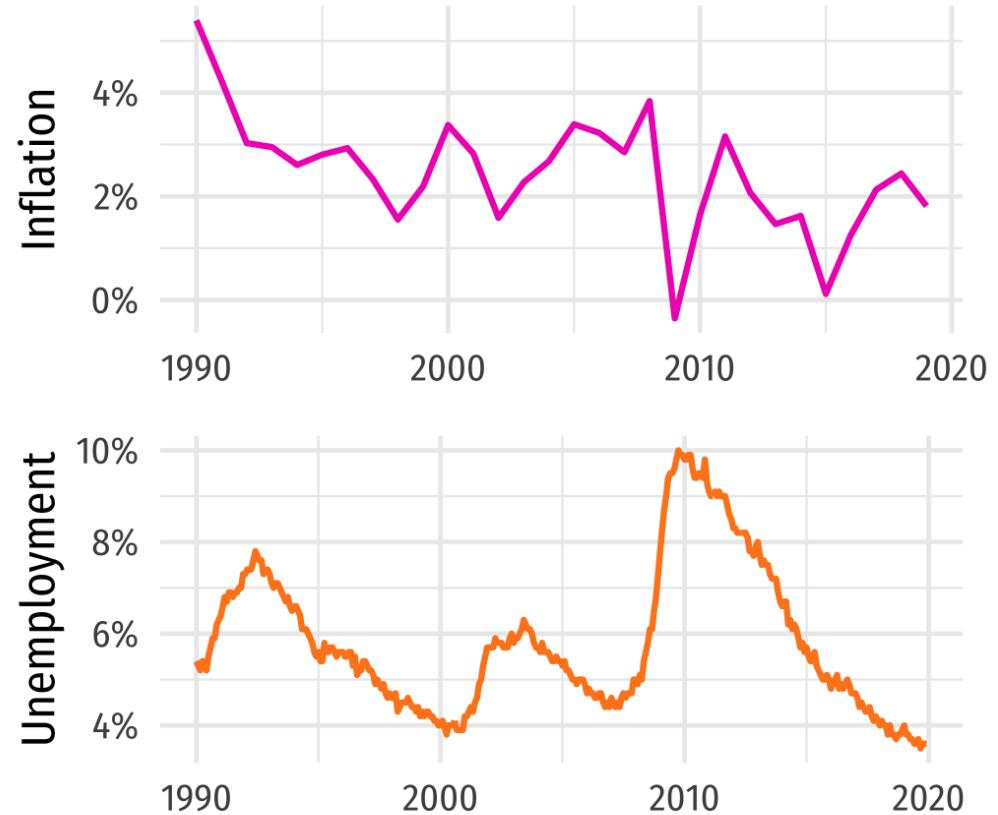
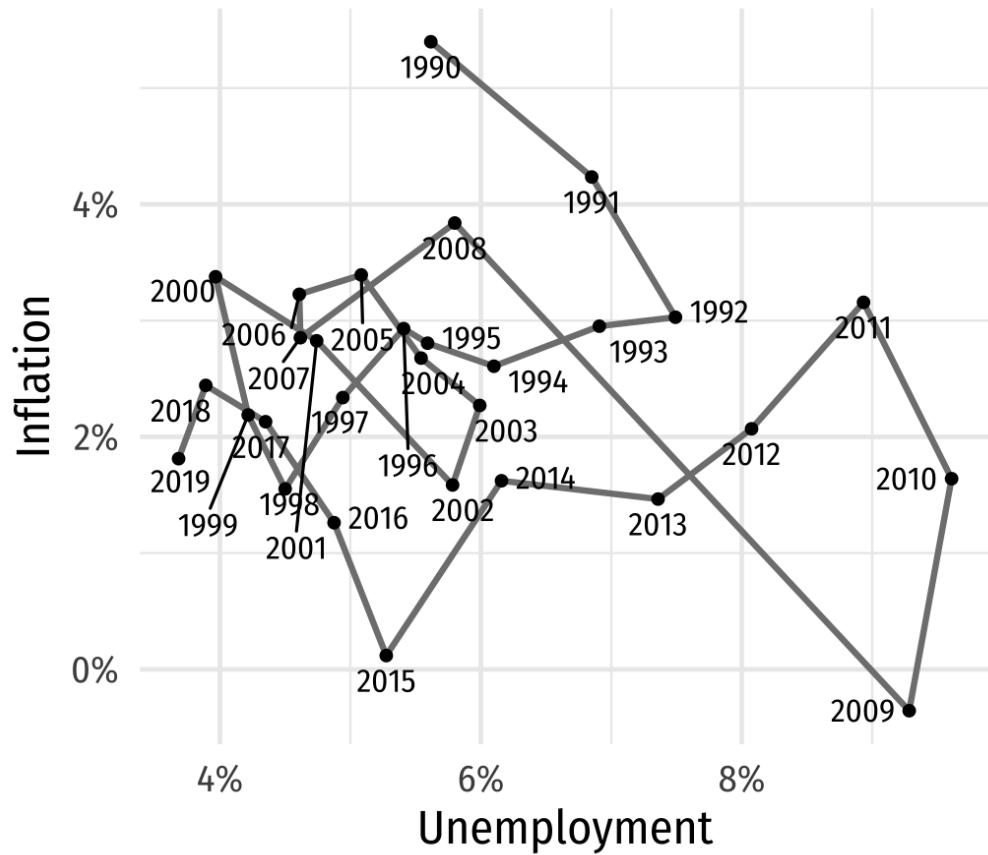
$$\log(Q) = \alpha + \beta \log(P) + \varepsilon$$

We could use regression to estimate  $\beta$

How might we interpret  $\beta$ ?

**But even here, a regular (not connected) scatter plot would be better!**

# Often it's better to use multiple plots



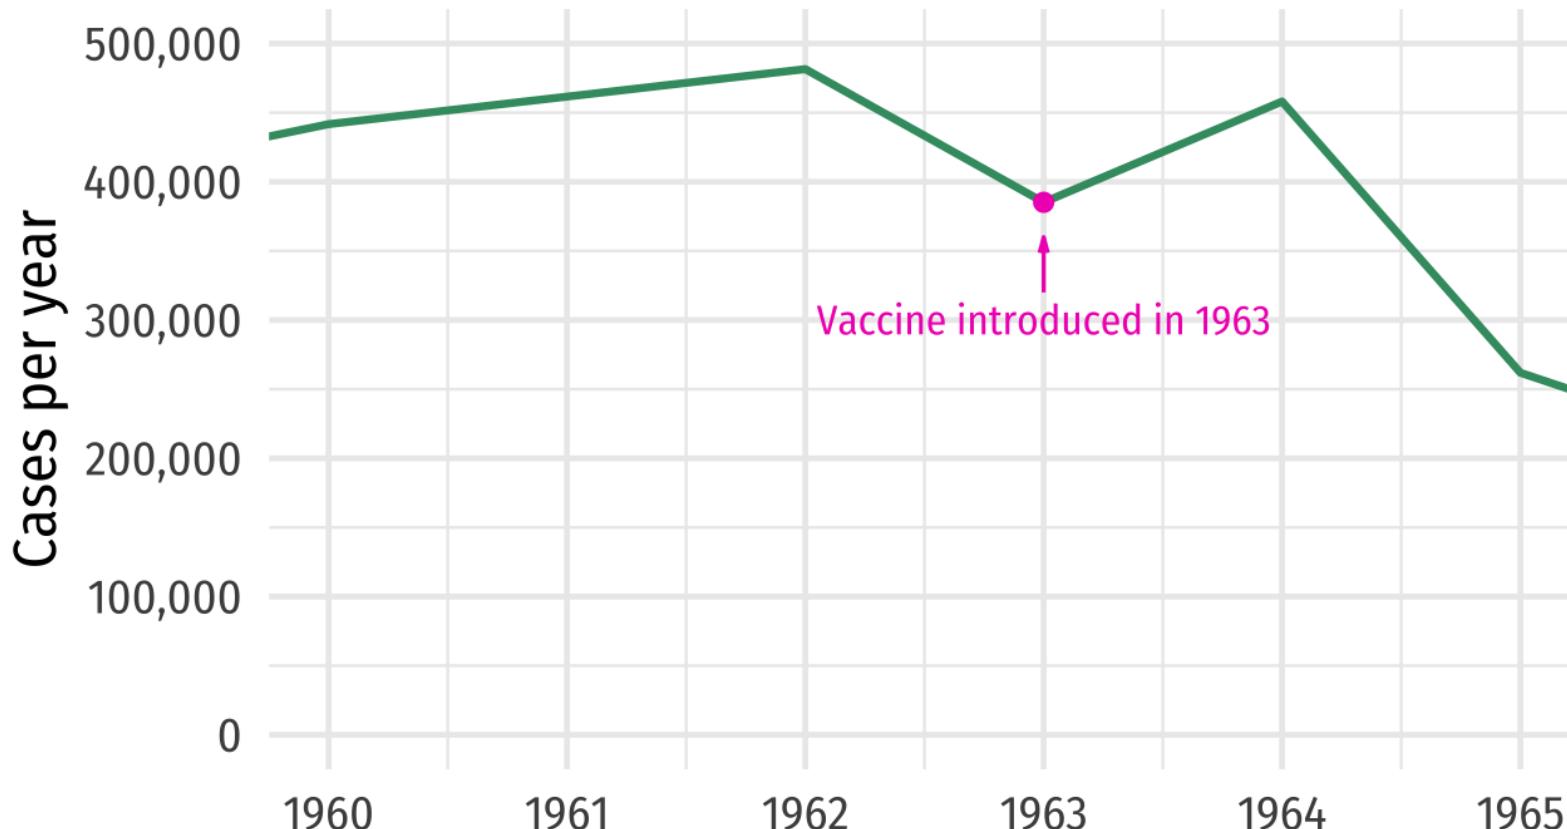
# Starting, ending, and decomposing time

# You have to start (and end) somewhere

You always have to choose start and end points

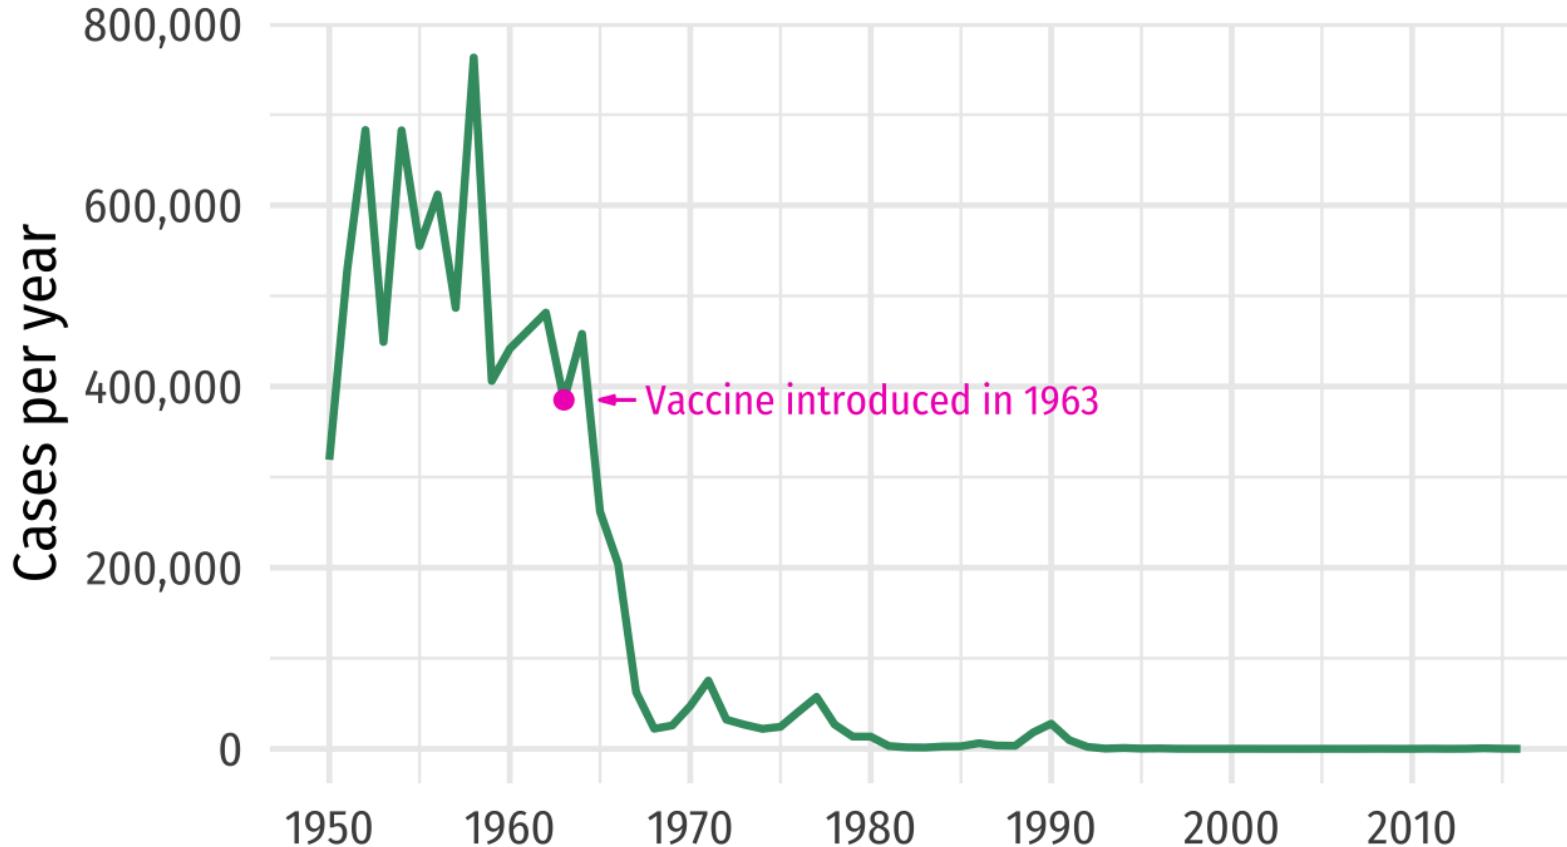
Start and end at reasonable times that help maintain the context of your story

# Measles vaccine was pretty effective



Source: CDC, Epidemiology and Prevention of Vaccine-Preventable Diseases, 13th Edition

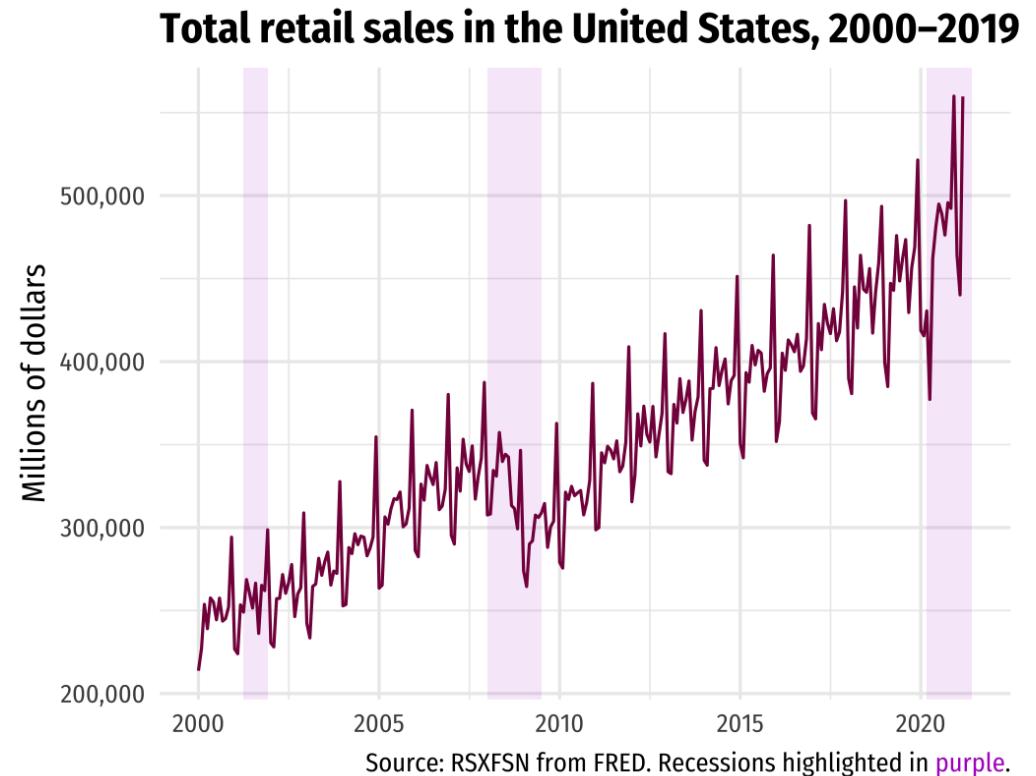
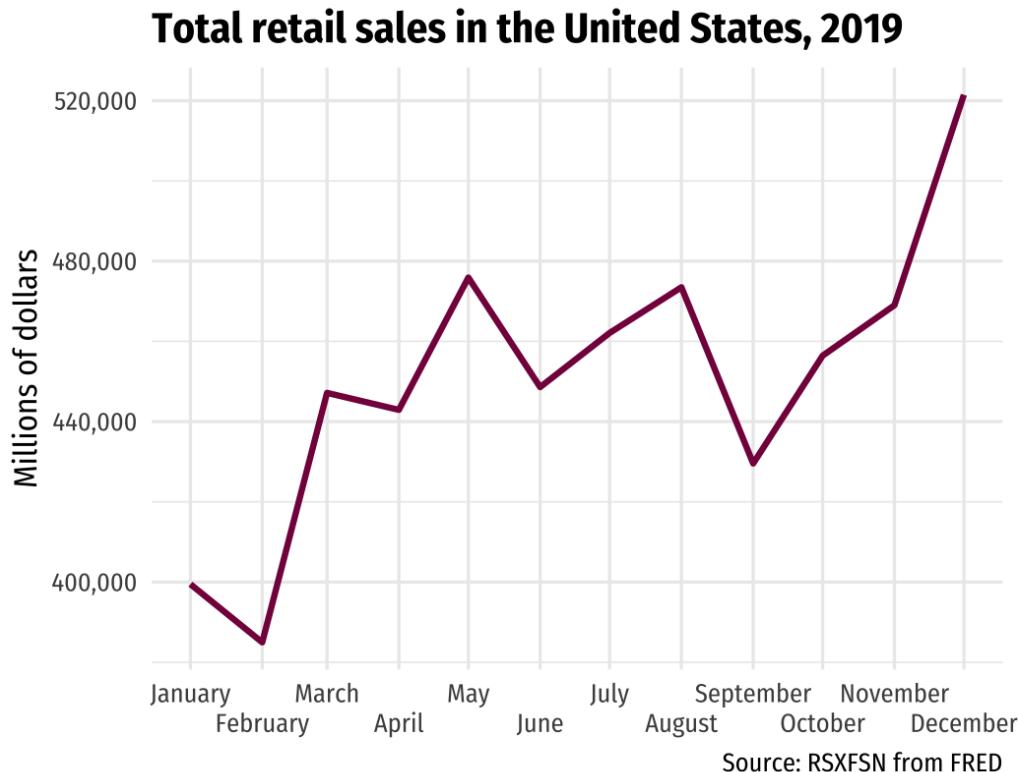
# Measles vaccine was *incredible*!



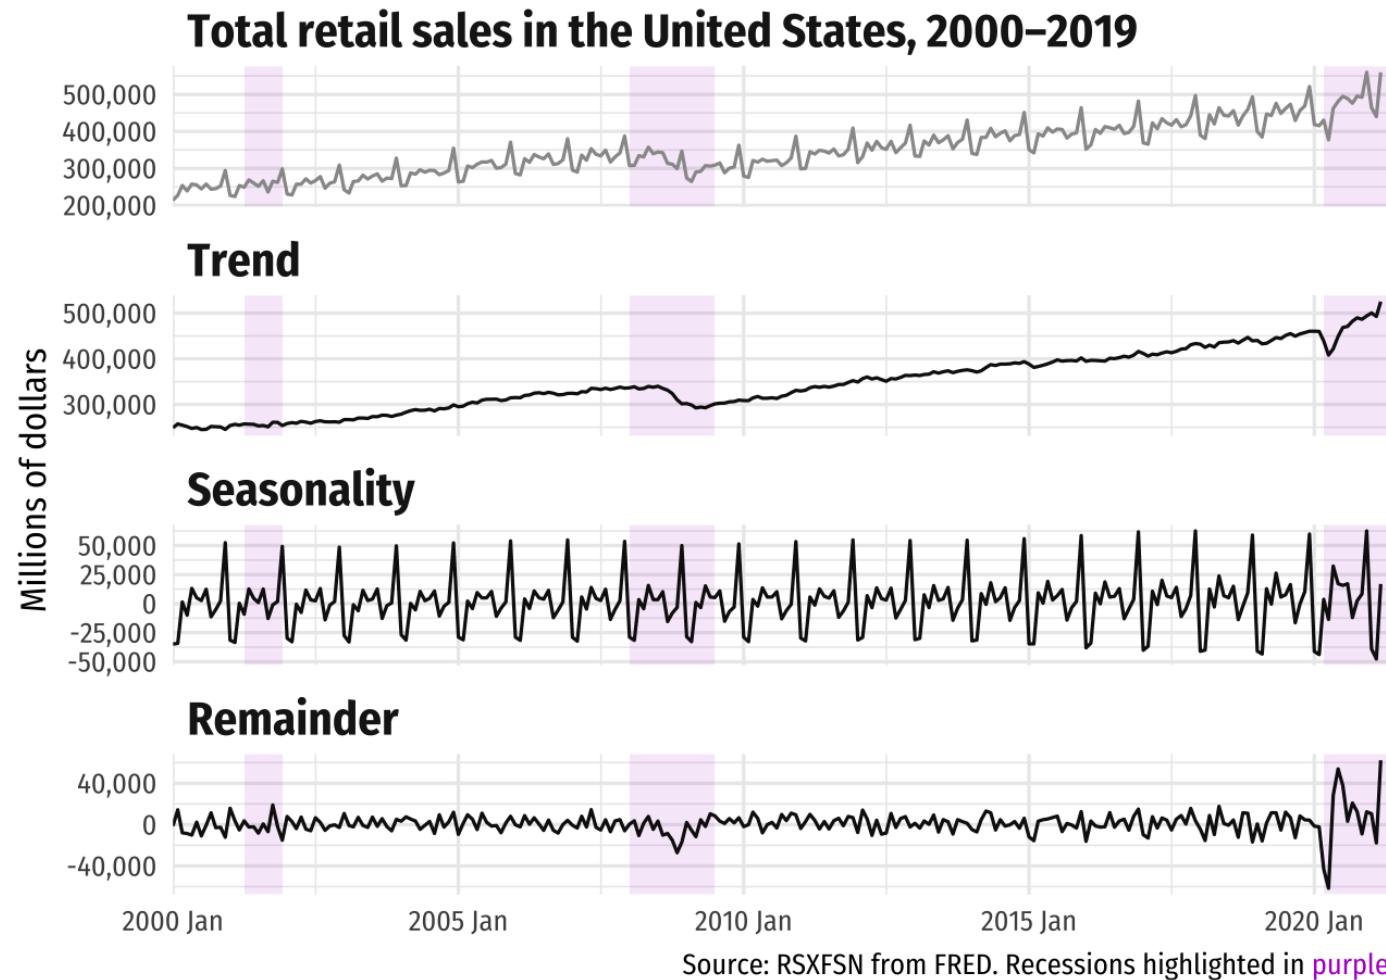
Source: CDC, Epidemiology and Prevention of  
Vaccine-Preventable Diseases, 13th Edition

# Seasonality

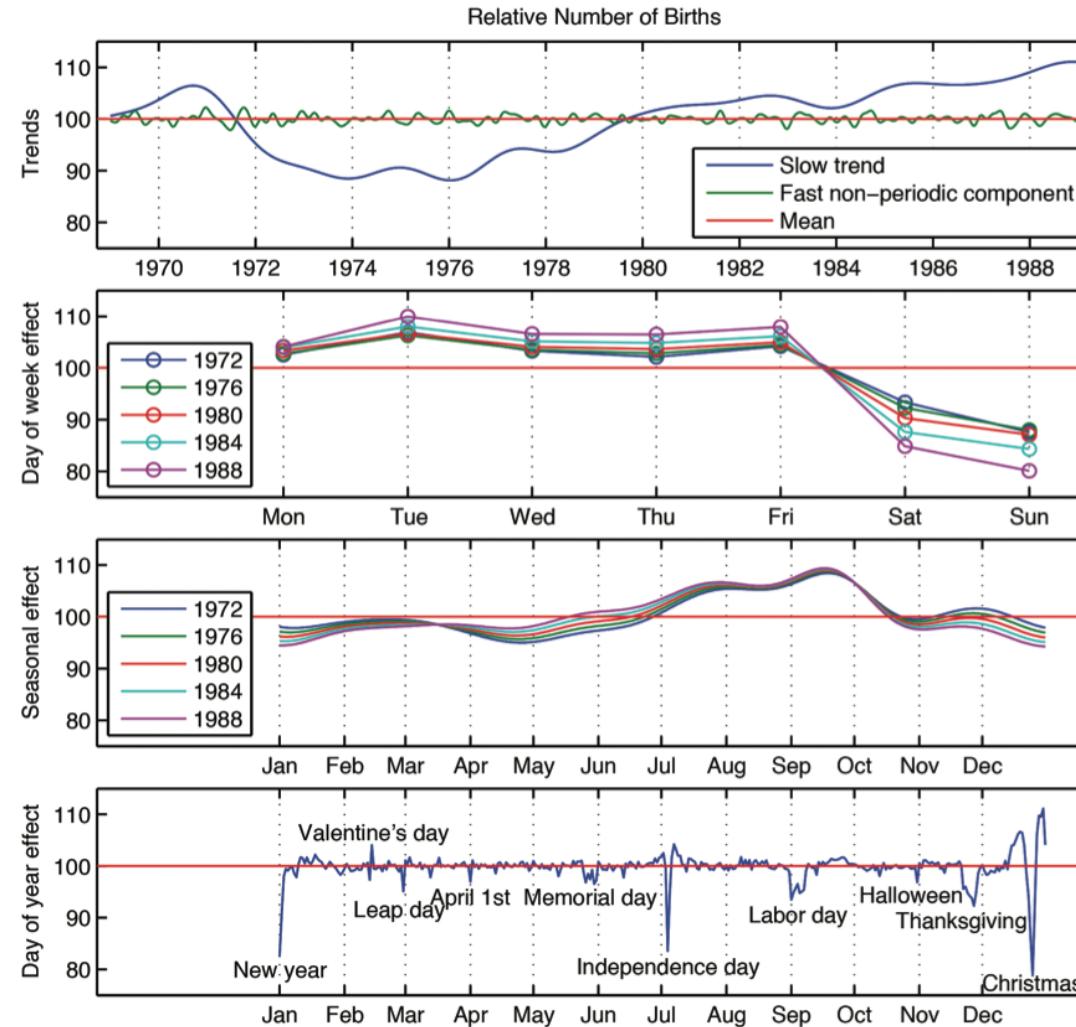
Don't mistake seasonality for actual trends



# Seasonal adjustment



# Birthday decomposition



# Dates and times in R

# Are dates and times special?

We made a bunch of plots without really thinking about data types

We can map "dates" to x regardless of whether they are `date` or `numeric` objects

```
class(gme_prices$date)
```

```
## [1] "Date"
```

```
class(solar_data$date)
```

```
## [1] "Date"
```

```
class(measles$Year)
```

```
## [1] "numeric"
```

When do we need to pay attention to the details?

1. Converting from strings to dates
2. Getting components of date-time data
3. Computing time spans

# Dates and times in R

R has "native" classes for storing calendar dates and times

- For background, see `?Dates` and `?DateTimeClasses`

The `lubridate` package offers convenient tools for working with dates and times

- Loaded automatically as part of the core tidyverse since tidyverse 2.0.0
- See `vignette("lubridate")` and Chapter 18 of R4DS (2e)

# Converting from strings to dates

Use `lubridate::ymd` to parse dates with **year**, **month**, and **day** components

```
svb_failure <- ymd("2023-03-10")  
svb_failure
```

```
## [1] "2023-03-10"
```

```
class(svb_failure)
```

```
## [1] "Date"
```

```
class("2023-03-10")
```

```
## [1] "character"
```

# Converting from strings to dates

`lubridate` offers many other functions for parsing dates

For example, instead of `ymd` we could use `mdy`:

```
signature_failure <- mdy("March 12, 2023")
```

```
signature_failure
```

```
## [1] "2023-03-12"
```

```
class(signature_failure)
```

```
## [1] "Date"
```

# Getting components of date-time data

What year did Silicon Valley Bank fail?

```
year(svb_failure)
```

```
## [1] 2023
```

What day of the week was that?

```
wday(svb_failure)
```

```
## [1] 6
```

```
wday(svb_failure, label = TRUE)
```

```
## [1] Fri  
## Levels: Sun < Mon < Tue < Wed < Thu < Fri < Sat
```

What month was that?

```
month(svb_failure)
```

```
## [1] 3
```

```
month(svb_failure, label = TRUE)
```

```
## [1] Mar  
## 12 Levels: Jan < Feb < Mar < Apr < May < Jun < Jul
```

# Computing time spans

How many days passed between independence and emancipation?

```
signature_failure - svb_failure
```

```
## Time difference of 2 days
```

How many days are left until classes end?

```
ymd(20230509) - today()
```

```
## Time difference of 51 days
```

These are **durations**

See Chapter 18 of R4DS (2e) for functions to compute **periods** and **intervals**